

Tutorial 1

CS5242: Neural Networks and Deep Learning

Zheng Zangwei

National University of Singapore

Agenda

1. Differentiation
2. PyTorch
3. Gradient Descend
4. Q&A Session

Differentiation

Recap: Importance of Gradient

Gradient is a key element in optimization.

In Deep learning, gradient descent is used for optimization. We must know the gradient of loss over all parameters (scalar-by-matrix differentiation).

Recap: Differentiation types

Types of matrix derivative

| Types | Scalar | Vector | Matrix |
|--------|--|---|--|
| Scalar | $\frac{\partial y}{\partial x}$ | $\frac{\partial \mathbf{y}}{\partial x}$ | $\frac{\partial \mathbf{Y}}{\partial x}$ |
| Vector | $\frac{\partial y}{\partial \mathbf{x}}$ | $\frac{\partial \mathbf{y}}{\partial \mathbf{x}}$ | |
| Matrix | $\frac{\partial y}{\partial \mathbf{X}}$ | | |

1. Gradient shape (shape check)
2. Gradient value (gradient rule)

Recap: Differentiation shapes

Our layout: **denominator layout**

This means the shape of our gradient is decided according to the shape of **denominator** and the transpose of **numerator**. ($\frac{\partial y}{\partial \mathbf{x}}$, then shape according to \mathbf{x} and \mathbf{y}^T)

Given scalar x , y , vector \mathbf{x} of shape $[n \times 1]$, vector \mathbf{y} of shape $[m \times 1]$, we have:

$$\frac{\partial y}{\partial x} : [1 \times 1] / [1 \times 1] \rightarrow [1 \times 1]$$

$$\frac{\partial \mathbf{y}}{\partial x} : [m \times 1] / [1 \times 1] \rightarrow [1 \times m]$$

$$\frac{\partial y}{\partial \mathbf{x}} : [1 \times 1] / [n \times 1] \rightarrow [n \times 1]$$

Recap: Differentiation shapes (Cont.)

Our layout: **denominator layout**

Given scalar x , y , vector \mathbf{x} of shape $[n \times 1]$, vector \mathbf{y} of shape $[m \times 1]$, matrix \mathbf{X} , \mathbf{Y} of shape $[m \times n]$

we have:

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} : [m \times 1] / [n \times 1] \rightarrow [n \times m]$$

$$\frac{\partial y}{\partial \mathbf{X}} : [1 \times 1] / [m \times n] \rightarrow [m \times n]$$

We only use enumerator layout for the following case:

$$\frac{\partial Y}{\partial x} : [m \times n] / [1 \times 1] \rightarrow ?$$

Recap: Differentiation rules

Ways to differentiate function with vectors:

1. By definition.
2. Use rules (Refer to Wiki and Matrix Cookbook).
3. Pushforward.

Recap: Differentiation rules (Cont.)

By definition

$$\frac{\partial y}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y}{\partial x_1} \\ \frac{\partial y}{\partial x_2} \\ \vdots \\ \frac{\partial y}{\partial x_n} \end{bmatrix}.$$

$$\frac{\partial \mathbf{y}}{\partial x} = \begin{bmatrix} \frac{\partial y_1}{\partial x} & \frac{\partial y_2}{\partial x} & \dots & \frac{\partial y_m}{\partial x} \end{bmatrix}.$$

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_2}{\partial x_1} & \dots & \frac{\partial y_m}{\partial x_1} \\ \frac{\partial y_1}{\partial x_2} & \frac{\partial y_2}{\partial x_2} & \dots & \frac{\partial y_m}{\partial x_2} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial y_1}{\partial x_n} & \frac{\partial y_2}{\partial x_n} & \dots & \frac{\partial y_m}{\partial x_n} \end{bmatrix}.$$

$$\frac{\partial y}{\partial \mathbf{X}} = \begin{bmatrix} \frac{\partial y}{\partial x_{11}} & \frac{\partial y}{\partial x_{12}} & \dots & \frac{\partial y}{\partial x_{1q}} \\ \frac{\partial y}{\partial x_{21}} & \frac{\partial y}{\partial x_{22}} & \dots & \frac{\partial y}{\partial x_{2q}} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial y}{\partial x_{p1}} & \frac{\partial y}{\partial x_{p2}} & \dots & \frac{\partial y}{\partial x_{pq}} \end{bmatrix}.$$

Recap: Differentiation rules (Cont.)

Use rules

| | | | | |
|-----------------------------------|-----|-----|-------|---------|
| $y =$ | a | x | Ax | $x^T A$ |
| $\frac{\partial y}{\partial x} =$ | ? | ? | A^T | ? |

| | | | | | |
|-----------------------------------|--------------------|---|---------------------------------|--------------------|----------------------|
| $y =$ | au $u = u(x)$ | vu $v = v(x), u = u(x)$ | $v + u$ $v = v(x), u = u(x)$ | Au $u = u(x)$ | $g(u)$ $u = u(x)$ |
| $\frac{\partial y}{\partial x} =$ | ? | $v \frac{\partial u}{\partial x} + \frac{\partial v}{\partial x} u^T$ | ? | ? | ? |

| | | | | |
|-----------------------------------|-----|---------------------------------|----------------------|-----------|
| $y =$ | a | $u^T v$ $v = v(x), u = u(x)$ | $g(u)$ $u = u(x)$ | $x^T A x$ |
| $\frac{\partial y}{\partial x} =$ | ? | ? | ? | ? |

Recap: Differentiation rules (Cont.)

Why rules hold? Use definition to prove.

Let's take $\frac{\partial \mathbf{Ax}}{\partial \mathbf{x}}$ as an example.

Shape check: \mathbf{x} [n x 1], \mathbf{A} [m x n]; \mathbf{Ax} [m x 1]

So the gradient \mathbf{g} is of shape [n x m].

Consider one position:

$$\begin{aligned} g_{i,j} &= \frac{\partial (\mathbf{Ax})_j}{\partial x_i} \\ &= \frac{\sum_{k=1}^n (A_{j,k} x_k)}{\partial x_i} \\ &= A_{j,i} \end{aligned}$$

So the gradient $\frac{\partial \mathbf{Ax}}{\partial \mathbf{x}} = \mathbf{A}^T$

Homework discussion

- Sigmoid function

$$f(x) = \frac{1}{1 + e^{-x}}$$

- Softmax

$$f(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}, 1 \leq i \leq n$$

- Softplus activation

$$f(x) = \frac{1}{\beta} \cdot \ln(1 + e^{\beta x})$$

Homework discussion (Cont.)

-

$$f(\mathbf{x}) = \mathbf{x}^T(\mathbf{A}\mathbf{x} + \mathbf{z})$$

- L2 loss

$$L(\mathbf{w}) = \frac{1}{2}(\mathbf{w}^T \mathbf{x} - y)^2$$

- L2 loss (multiple examples)

$$L(\mathbf{w}) = \frac{1}{2m} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

-

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}, L = \|\mathbf{z} - \mathbf{y}\|^2$$

Identities

PyTorch

How does program do differentiation?

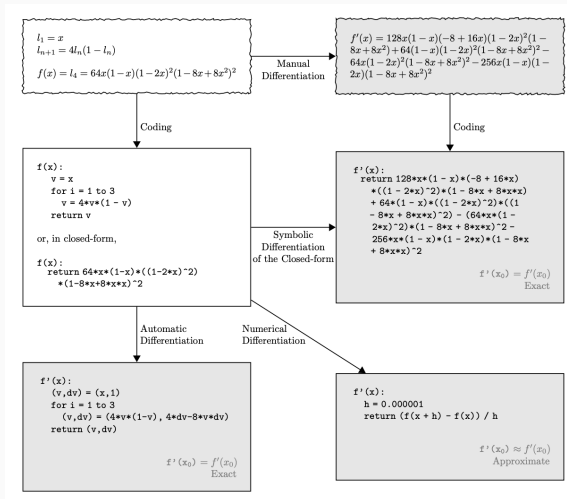


Figure 1: Source: Baydin, Atılım Gunes, et al. "Automatic differentiation in machine learning: a survey." *Journal of Machine Learning Research* 18 (2018): 1-43.

How does program do differentiation?

- Manual Differentiation
- Symbolic Differentiation of the Closed-form (e.g., Mathematica, Maple)
- Numerical Differentiation
- Automatic Differentiation (PyTorch, Tensorflow)

How does program do differentiation? (Cont.)

| Forward Primal Trace | Forward Tangent (Derivative) Trace |
|--|---|
| $v_{-1} = x_1 = 2$ | $\dot{v}_{-1} = \dot{x}_1 = 1$ |
| $v_0 = x_2 = 5$ | $\dot{v}_0 = \dot{x}_2 = 0$ |
| $v_1 = \ln v_{-1} = \ln 2$ | $\dot{v}_1 = \dot{v}_{-1}/v_{-1} = 1/2$ |
| $v_2 = v_{-1} \times v_0 = 2 \times 5$ | $\dot{v}_2 = \dot{v}_{-1} \times v_0 + \dot{v}_0 \times v_{-1} = 1 \times 5 + 0 \times 2$ |
| $v_3 = \sin v_0 = \sin 5$ | $\dot{v}_3 = \dot{v}_0 \times \cos v_0 = 0 \times \cos 5$ |
| $v_4 = v_1 + v_2 = 0.693 + 10$ | $\dot{v}_4 = \dot{v}_1 + \dot{v}_2 = 0.5 + 5$ |
| $v_5 = v_4 - v_3 = 10.693 + 0.959$ | $\dot{v}_5 = \dot{v}_4 - \dot{v}_3 = 5.5 - 0$ |
| $y = v_5 = 11.652$ | $\dot{y} = \dot{v}_5 = 5.5$ |

Figure 2: Source: Baydin, Atilim Gunes, et al. "Automatic differentiation in machine learning: a survey." *Journal of Machine Learning Research* 18 (2018): 1-43.

1. PyTorch
2. Colab
3. Homework

Gradient Descend

Training by gradient descent

- Gradient descent algorithm for optimization
- Set $w = 0.1$ or a random number
- Repeat
 - For each data sample, compute $\tilde{y} = xw + b$
 - Compute the average loss, $\sum_{\langle x, y \rangle \in S_{train}} L(x, y | w, b) / |S_{train}|$
 - Compute $\frac{\partial J}{\partial w}$
 - Update $w = w - \alpha \frac{\partial J}{\partial w}$

Update w, b repeatedly...

Learning rate

A good learning rate α make sure the process converges and converges fast.

For a simple case,

- Converge gradually
- Oscillate but converge
- Diverge

In the practice, more complicated.

- local minimal
- flatness
-

So learning rate is often the most important hyper-parameter to tune in DL.

Homework discussion

Consider a linear regression without intercept $y = xw$, $x \in \mathbb{R}$, $w \in \mathbb{R}$. L2 loss and gradient descent are used. Initial $w = 0$ and learning rate is α . Suppose we only have one example $x = 1, y = 100$ (which is not a setting in the reality and we use this toy example for ease of computation).

1. Show how gradient descent works for $\alpha = 0.5, 1.5, 2.5$.
2. Give the condition of α that gradient descent starts to oscillate around the optimal position. Give the condition of α that gradient descent can converge.
3. Try to prove your statement in 2). (Hint: consider $|100 - w_t|$)

Q&A Session
