

# 4. Convolutional Neural Network (CNN) Basics

CS 5242 Neural Networks and Deep Learning

YOU, Yang

06.09.2022

# Recap

- Backpropagation
  - A modular way to compute the gradients of the loss w.r.t each parameter
  - Represent the computation using a graph
    - A node for each operation, and an edge for each variable
    - Each operation implements two functions: forward and backward
  - Apply chain rules against the graph
    - Forward the data through every node in topological order
    - Backward the gradient through every node in the reverse order

# Recap

- Mini-batch stochastic gradient descent (SGD)
  - Reduces the chance of local optimal points and saddle points from GD
  - More stable than standard SGD
  - Extensions: Momentum, RMSProp, Adam
    - Exploiting historical updates
    - Adaptive learning rate per parameter
- Training tricks
  - Parameter initialization
  - Data normalization
  - Regularization:
    - Early stopping

# Agenda

- Convolution basics
- 1D and 2D convolution operations
- Pooling operations

# Convolutional Neural Networks

- Referred to in short form as ConvNets or CNNs
- Most frequently used for image(-related data):
  - Image classification
  - Object detection
  - (Medical) image segmentation
  - [Face recognition](#)
  - [Image generation](#)
  - [Art composition](#)

# From MLP to CNN

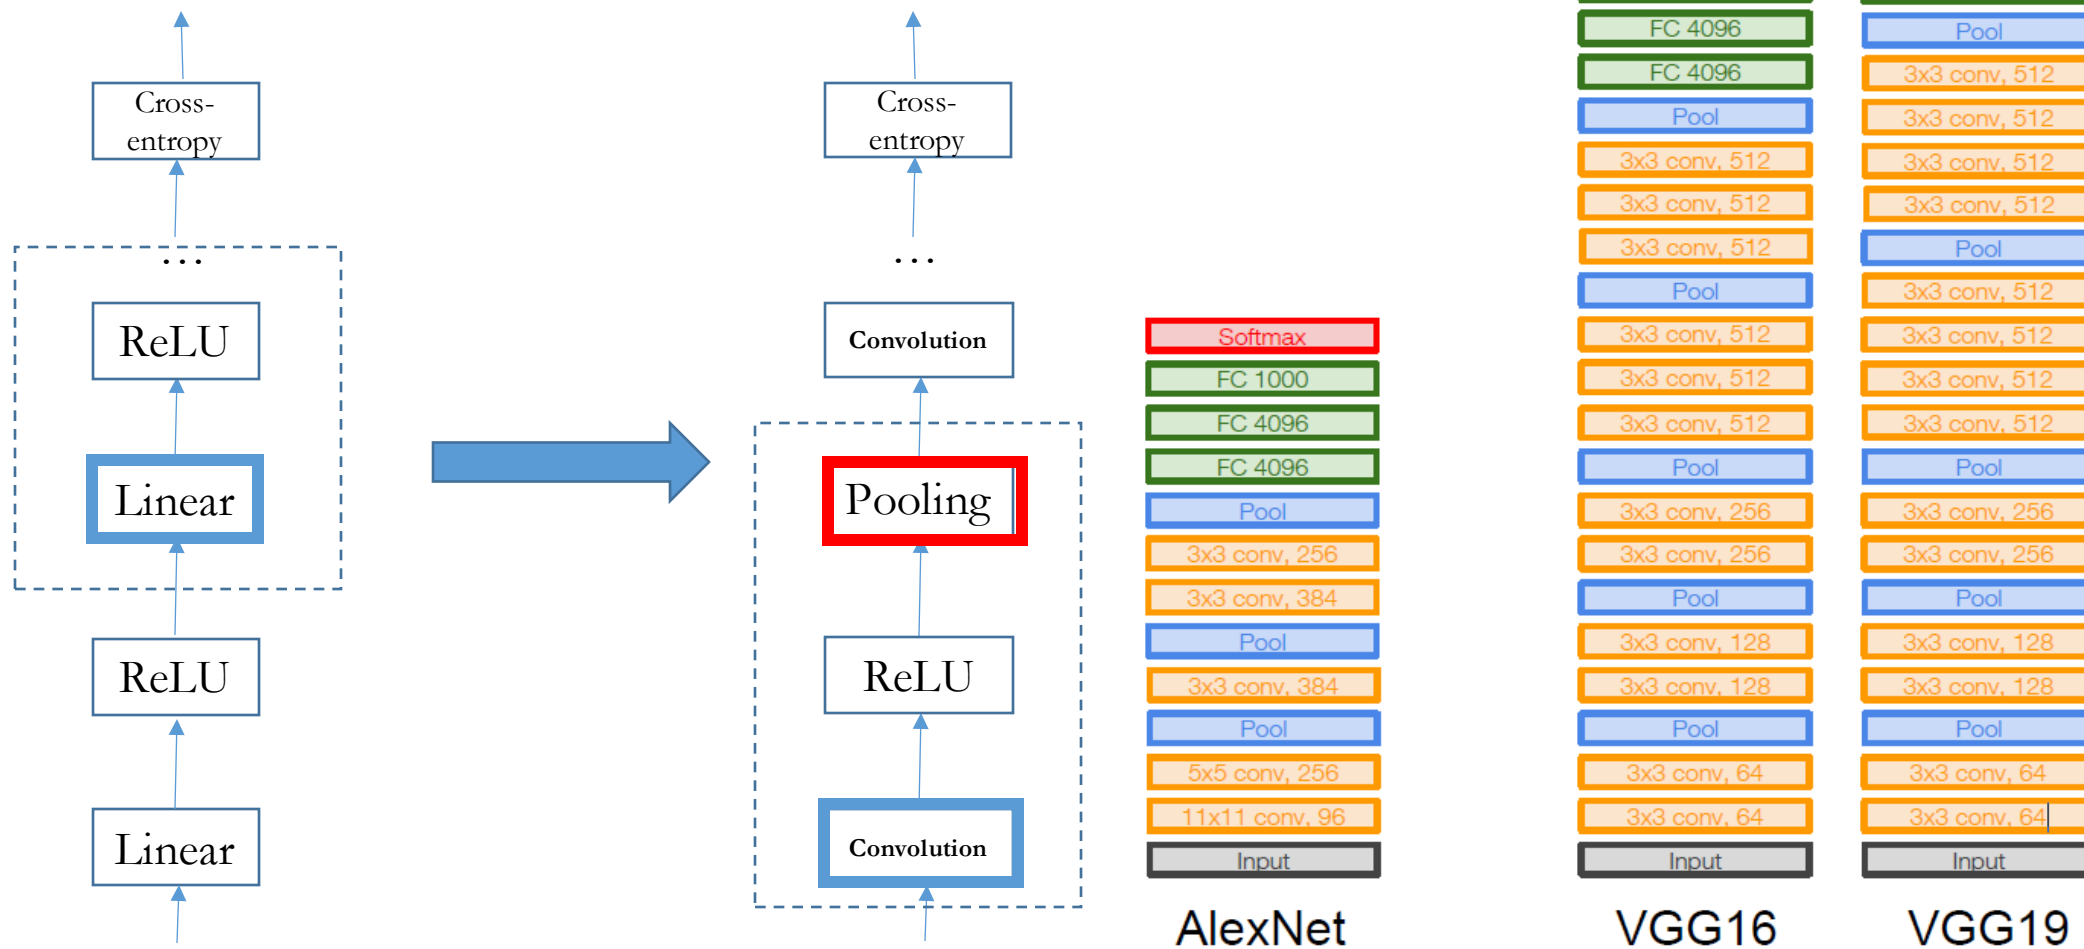


Image source: Stanford cs231n

# Convolution

- A linear transformation
- Since output is a “feature” of the input, convolution can be considered feature extraction
- 1D / 2D / 3D
  - 1D: Text processing
  - 2D: Image processing
  - 3D: 3D data, CT, microscopy, etc.
- mD convolution
  - “mD” comes from “m” dimensions of the source data
  - Can apply 1D convolution over 1D or 2D data;
  - Can apply 2D convolution over 2D or 3D data;

## Feature Visualization

How neural networks build up their understanding of images



[Feature visualization](#)



# 1D convolution

$$y_t = \sum_{i=0}^{k-1} w_i \times x_{t+i}$$

Cross-correlation operation [[link](#)]




- In CNNs, convolution refers to cross-correlation
- $\mathbf{W}$  is called the kernel/filter; length  $k$ 
  - "weights" or parameters to be trained;
- $\mathbf{x}$  is the input; length  $n$
- the applied or input area, i.e.  $t, t+1, \dots, t+k-1$  is called the receptive field
  - one receptive field generates one output value
- $y_t$  is the output feature; length  $o$

In signal processing, cross-correlation is a measure of similarity of two series as a function of the displacement of one relative to the other.

It is also known as a sliding dot product.



# 1D convolution



$$3 \times 1 + 2 \times 0 + 2 \times 2 = 7$$

$$3 \times 0 + 2 \times 2 + 2 \times 2 = 8$$

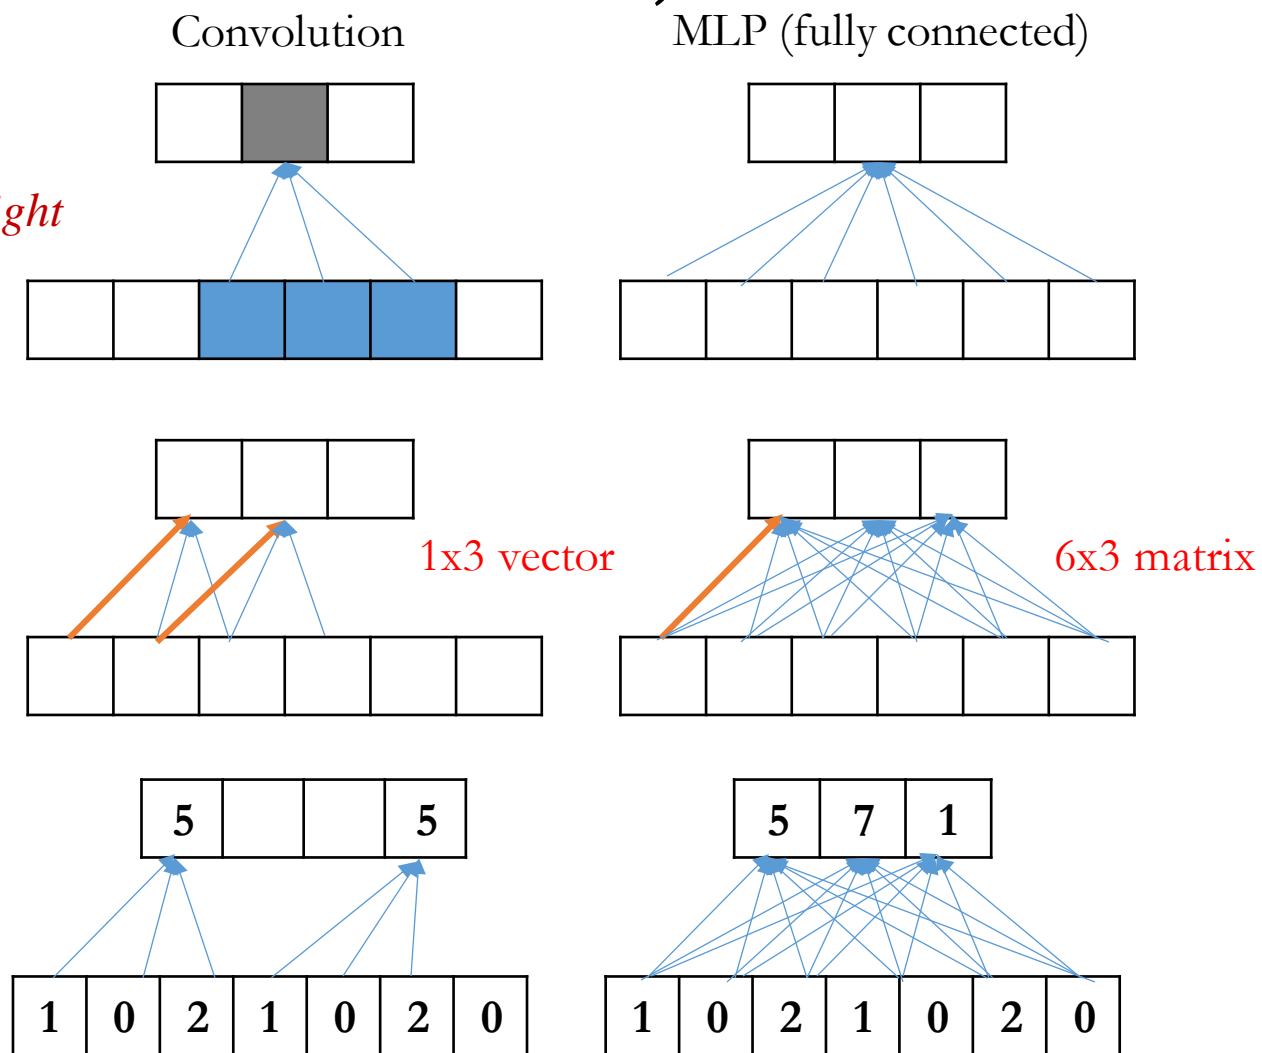
$$3 \times 2 + 2 \times 2 + 2 \times 3 = 16$$

$$3 \times 2 + 2 \times 2 + 2 \times 1 = 14$$

$$3 \times 3 + 2 \times 1 + 2 \times 0 = 11$$

# Properties (Why is convolution better?)

- Sparse connection:
  - each output is connected only to inputs within receptive field vs. all inputs
  - → fewer parameters (each output needs 3 vs. 6 in example)
  - → less overfitting
- Weight sharing vs. unique weights
  - Regularization
  - Less overfitting
- Location or Spatial invariant
  - Function transformations should not depend on the location within the image, i.e.
  - Make the same prediction no matter where the object is in the image



# Location or Spatial Invariant

- You can recognize an object even its appearance varies in some way
- Convolution operator commutes with respect to translation
  - If you convolve  $f$  with  $g$ , it doesn't matter if you translate the convolved output  $f * g$ , or you translate  $f$  or  $g$  first, then convolve them.
  - <https://en.wikipedia.org/wiki/Convolution>
- Location or Spatial invariant
  - Function transformations should not depend on the location within the image, i.e.
  - Make the same prediction no matter where the object is in the image

Translation Invariance



Rotation/Viewpoint Invariance



Size Invariance



Slide credit: Matt Krause

# Perceptron, MLP and Convolution

## Perceptron

Perceptron is too simple

- underfitting
- add more layers
- MLP

## MLP



MLP has too many parameters

- High dimension
- difficult to optimize and overfitting
- CNN (with more regularization)

## CNN

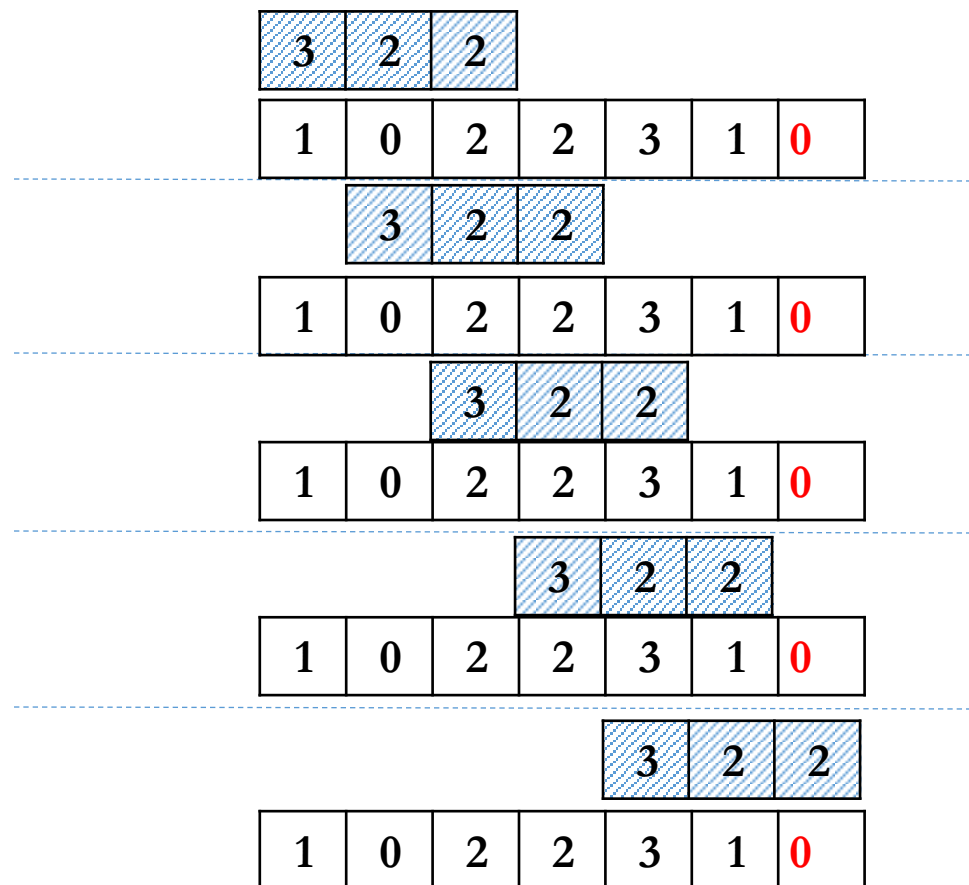
# Padding

- How to determine the edge values?  
Ignore non-valid regions?
- $o = n - k + 1$ 
  - $n$ : input length,  $k$ : kernel length,  $o$ : output length
  - Output is shorter than input
- To retain the resolution/size
  - Pad with extra values (usually 0s)



# Padding

- Manual padding ( $p$ )
  - Output feature values for  $p = 1$ 
    - $3 \times 1 + 2 \times 0 + 2 \times 2 = 7$
    - $3 \times 0 + 2 \times 2 + 2 \times 2 = 8$
    - $3 \times 2 + 2 \times 2 + 2 \times 3 = 16$
    - $3 \times 2 + 2 \times 2 + 2 \times 1 = 14$
    - $3 \times 3 + 2 \times 1 + 2 \times 0 = 11$
- What value to pad with?
  - usu.  $k \ll n$  so value doesn't matter too much; so don't bother tuning
  - 0 picked for convenience
- Operation supported in many deep learning libraries, e.g.
  - Torch, PyTorch, Caffe, [SINGA](#)





# Padding amount ( $p$ )

- Given padding ( $p$ ), what will the output length be?
- Kernel size/length:  $k$ , input length:  $n$
- Without padding:
  - # outputs  $o = n - k + 1$
- With padding:
  - # outputs  $o = (n + p) - k + 1$
- padding length ( $p$ ) can be set manually or automatically
- 2 special automatic settings:
  - consider only “valid” convolutions  $\rightarrow p = 0$
  - same length output as input  $\rightarrow$  “same”
    - $o = n \rightarrow p = f(k)$   $p = k - 1$

# “Valid” Convolution

- No padding ( $p = 0$ )
  - # inputs denoted as  $n$
  - # outputs  $o = n - k + 1 = 6 - 3 + 1 = 4$
  - Output feature values
    - $3 \times 1 + 2 \times 0 + 2 \times 2 = 7$
    - $3 \times 0 + 2 \times 2 + 2 \times 2 = 8$
    - $3 \times 2 + 2 \times 2 + 2 \times 3 = 16$
    - $3 \times 2 + 2 \times 2 + 2 \times 1 = 14$
  - Outputs become shorter
- Is an option to be set in library





# Same Padding

- Same padding ( $p$ ?)
  - $o = n = n + p - k + 1$
  - $p = k - 1$

• Left padding =  $\lfloor p/2 \rfloor$

• Right padding =  $\lfloor p/2 \rfloor$

If  $p$  is an odd number, libraries typically assign 1 less to the left than right or vice versa.







- Output values

- $3 \times 0 + 2 \times 1 + 2 \times 0 = 2$
- $3 \times 1 + 2 \times 0 + 2 \times 2 = 7$
- $3 \times 0 + 2 \times 2 + 2 \times 2 = 8$
- $3 \times 2 + 2 \times 2 + 2 \times 3 = 16$
- $3 \times 2 + 2 \times 3 + 2 \times 1 = 14$
- $3 \times 3 + 2 \times 1 + 2 \times 0 = 11$

For  $n$  an integer,  $\lfloor n \rfloor = \lceil n \rceil = [n] = n$ .



Examples

$x$	Floor $\lfloor x \rfloor$	Ceiling $\lceil x \rceil$	Fractional part $\{x\}$
2	2	2	0
2.4	2	3	0.4
2.9	2	3	0.9
-2.7	-3	-2	0.3
-2	-2	-2	0





# Stride

- How many steps to move towards the next receptive field
  - so far, we have considered only a stride of 1, where the kernel is applied directly to the next element in the array
  - $s > 1$ , skip some elements
    - Faster to compute
    - *Effective receptive field size increases quickly*



# Stride

- Increasing the stride is computationally faster, since we compute less convolutions
- Resulting output with higher stride is subsequently shorter
- $$o = \left\lfloor \frac{n+p-k}{s} \right\rfloor + 1$$
- Tunable hyperparameter, which may also vary depending on the layer



# Stride

- Exact matching
  - With padding  $p$  ( $=1$ )
  - $o = \left\lfloor \frac{n+p-k}{s} \right\rfloor + 1$
  - $(6+1-3)/2+1=3$

3	2	1			
1	0	2	2	3	1

3	2	2				
1	0	2	2	3	1	0

		<div><div>3</div><div>2</div><div>2</div></div>				
1	0	2	2	3	1	0

				3	2	2
1	0	2	2	3	1	0

# Stride

- Not exact matching
  - With padding  $p$  ( $=2$ )
  - $o = \left\lfloor \frac{n+p-k}{s} \right\rfloor + 1$
  - $(6+2-3)/2+1=3$

This last computation is not valid.  
This equation works in both situations.

3	2	1			
1	0	2	2	3	1

3	2	2					
1	0	2	2	3	1	0	0

								<div>322</div>										
1	0	2	2	3	1	0	0											





				3	2	2	
1	0	2	2	3	1	0	0

						3	2	2
1	0	2	2	3	1	0	0	



# Stride (for Tensorflow)

- When stride  $> 1$ 
  - Valid padding,  $p = 0$
  - Same padding, ?
- output length cannot be equal to the input, since a stride greater than 1 will shorten the output
- “same” is defined as:
  - $o = \left\lceil \frac{n}{s} \right\rceil, p = ?$   
ceiling operation.



# Stride (for Tensorflow)

- When stride > 1
  - Valid padding,  $p = 0$
  - Same padding,  $o = \left\lceil \frac{n}{s} \right\rceil$ ,  $p = ?$





$$o = \left\lceil \frac{n+p-k}{s} \right\rceil + 1$$

$$\frac{n+p-k}{s} \geq o - 1$$

$$p \geq s(o - 1) + k - n$$

$$p = \max(s(o - 1) + k - n, 0)$$

Tensorflow  
internally computes  
this p value when  
we set “same”  
padding for strides  
greater than 1.



# Summary by an Example (2-minute Quiz)

- Input Length = 13; Stride = 5; Kernel Length = 6
- "valid" method: no padding, drop the non-valid region

How many elements in input vector will be dropped?

- "same" method: padding at both ends

$$p = \max(s(o-1) + k - n, 0)$$

What is the padding length? Where?



## Summary by an Example

- Input Length = 13; Stride = 5; Kernel Length = 6
- "valid" method: no padding, drop the non-valid region

```
inputs:      1  2  3  4  5  6  7  8  9  10 11 (12 13)
             |_____|
                               |_____|
                               dropped
```

$$p = \max(s(o-1) + k - n, 0)$$

$$p = \max(s(\text{ceiling}(n/s)-1) + k - n, 0)$$

$$p = \max(5(\text{ceiling}(13/5)-1) + 6 - 13, 0)$$

$$p = \max(5(\text{ceiling}(2.6)-1) + 6 - 13, 0)$$

$$p = \max(5(3-1) + 6 - 13, 0)$$

$$p = \max(10 + 6 - 13, 0) = 3$$

- "same" method: padding at both ends

```

inputs:      pad|                                     |pad
              0 |1  2  3  4  5  6  7  8  9  10 11 12 13|0  0
              |_____|
                      |_____|
                              |_____|

```

# Computing

- Conv1D
  - Forward(x, w)

 $w^T$ 

3	2	1
---	---	---

 $x^T$ 

1	0	2	2	3	1
---	---	---	---	---	---

S=2, k=3, p=1

1	0	2	2	3	1	0
1	0	2	2	3	1	0
1	0	2	2	3	1	0

# Forward

 $w^T$ 

3	2	1
---	---	---

 $x^T$ 

1	0	2	2	3	1
---	---	---	---	---	---

$S=2, k=3, p=1$

Convolution via dot products in a for loop.

```
x_pad ← x
for t in range(o):
    y[t] = dot(w, x_pad[t*s:t*s+k])
```

Vectorization (img2col)

More efficient vectorized implementation; convert each receptive field as a single column.

```
for t in range(o):
    X[:, t] = x_pad[t*s:t*s+k]
y = dot(w, X)
```

Try out in Google colab!

1	0	2	2	3	1	0
1	0	2	2	3	1	0
1	0	2	2	3	1	0



receptive field to column

 $w^T$ 

3	2	1
---	---	---

1	2	3
0	2	1
2	3	0

 $X$ 


$dot()$

The result  $y$  is a row vector; if we compute  $y$  as  $X^T w$ , then the result  $y$  is a column vector

Converting many dot products to a vector-matrix multiplication

# 2D Convolution

- 2D convolution follows the same principles, but the inputs, kernels and outputs are generalized into 2D matrices instead of 1D vectors




$$Y_{i,j} = \sum_{a=0}^{k_h-1} \sum_{b=0}^{k_w-1} X_{i*s_h+a, j*s_w+b} \times W_{a,b}$$

Is this a matrix-matrix multiplication?

# 2D Convolution

- 2D convolution follows the same principles, but the inputs, kernels and outputs are generalized into 2D matrices instead of 1D vectors




$$Y_{i,j} = \sum_{a=0}^{k_h-1} \sum_{b=0}^{k_w-1} X_{i*s_h+a, j*s_w+b} \times W_{a,b}$$

Is this a matrix-matrix multiplication?

Element-wise operation:  $O(n^2)$  not  $O(n^3)$

# 2D Convolution

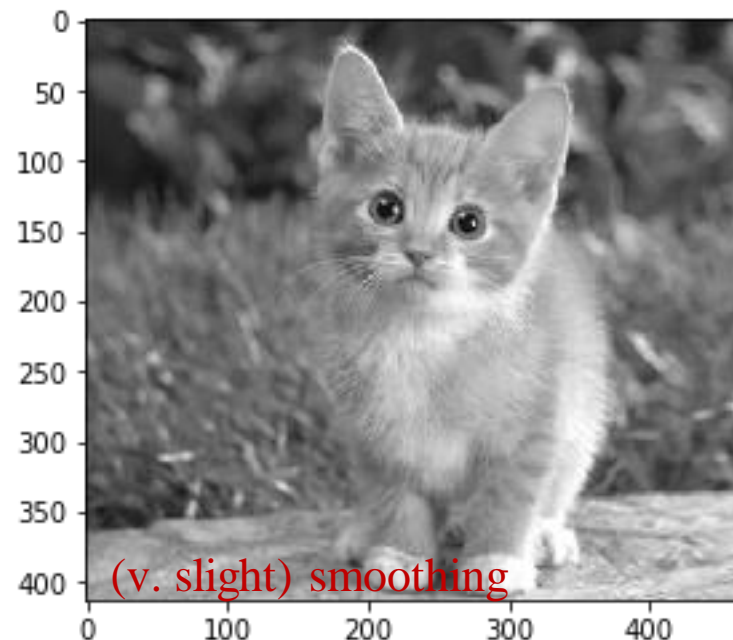






Source: [http://deeplearning.net/software/theano/tutorial/conv\\_arithmetic.html](http://deeplearning.net/software/theano/tutorial/conv_arithmetic.html)

# 2D Convolution

In deep learning, we learn the kernels or weights which gives us good predictions e.g. classification, regression etc.

<http://setosa.io/ev/image-kernels/>



We get different feature maps depending on the kernel used.

# 2D Convolution

$$\bullet Y_{i,j} = \sum_{a=0}^{k_h-1} \sum_{b=0}^{k_w-1} X_{i*s_h+a, j*s_w+b} \times W_{a,b}$$



1	2	3	2
2	3	2	5
1	2	2	1
0	2	1	2

1	2	3	2
2	3	2	5
1	2	2	1
0	2	1	2

-1	1
-1	1



2		

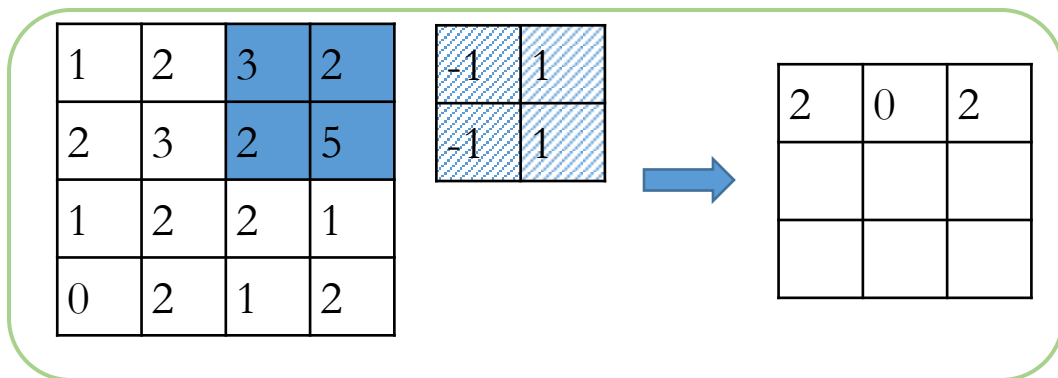




# 2D Convolution




1	2	3	2
2	3	2	5
1	2	2	1
0	2	1	2

$$\bullet Y_{i,j} = \sum_{a=0}^{k_h-1} \sum_{b=0}^{k_w-1} X_{i*s_h+a, j*s_w+b} \times W_{a,b}$$



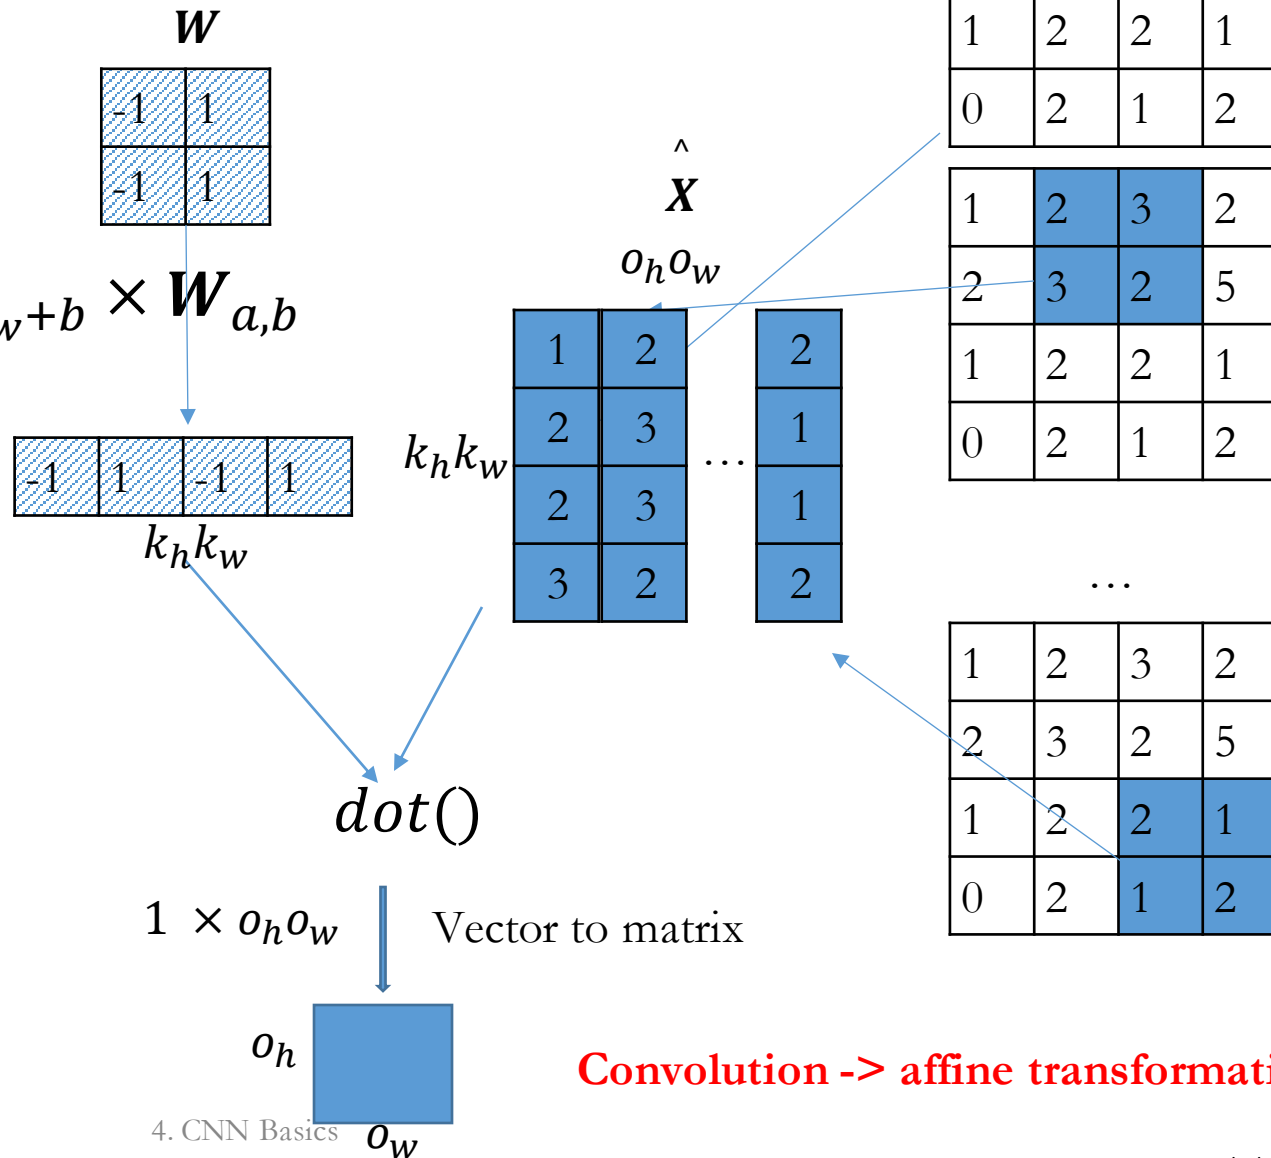
...



## Implementation

- Img2Col
  - Convert each receptive field into a column

- $\mathbf{Y}_{i,j} = \sum_{a=0}^{k_h-1} \sum_{b=0}^{k_w-1} \mathbf{X}_{i*s_h+a, j*s_w+b} \times \mathbf{W}_{a,b}$




## Convolution -> affine transformation

# Statistics


- Shape:
  - Input  $\mathbf{X} \in R^{n_h \times n_w}$
  - Kernel  $\mathbf{W} \in R^{k_h \times k_w}$
  - Output  $\mathbf{Y} \in R^{o_h \times o_w}$
- Parameter size
  - $k_h \times k_w$
- Output shape
  - $(o_h, o_w) = (\left\lfloor \frac{n_h + p_h - k_h}{s_h} \right\rfloor + 1, \left\lfloor \frac{n_w + p_w - k_w}{s_w} \right\rfloor + 1)$
- Computation cost
  - $O(k_h \times k_w \times o_h \times o_w)$  (float multiplication ops, FLOP)

# 2D Convolution

- Multiple kernels/filters



-1	1
-1	1



1	1
1	1



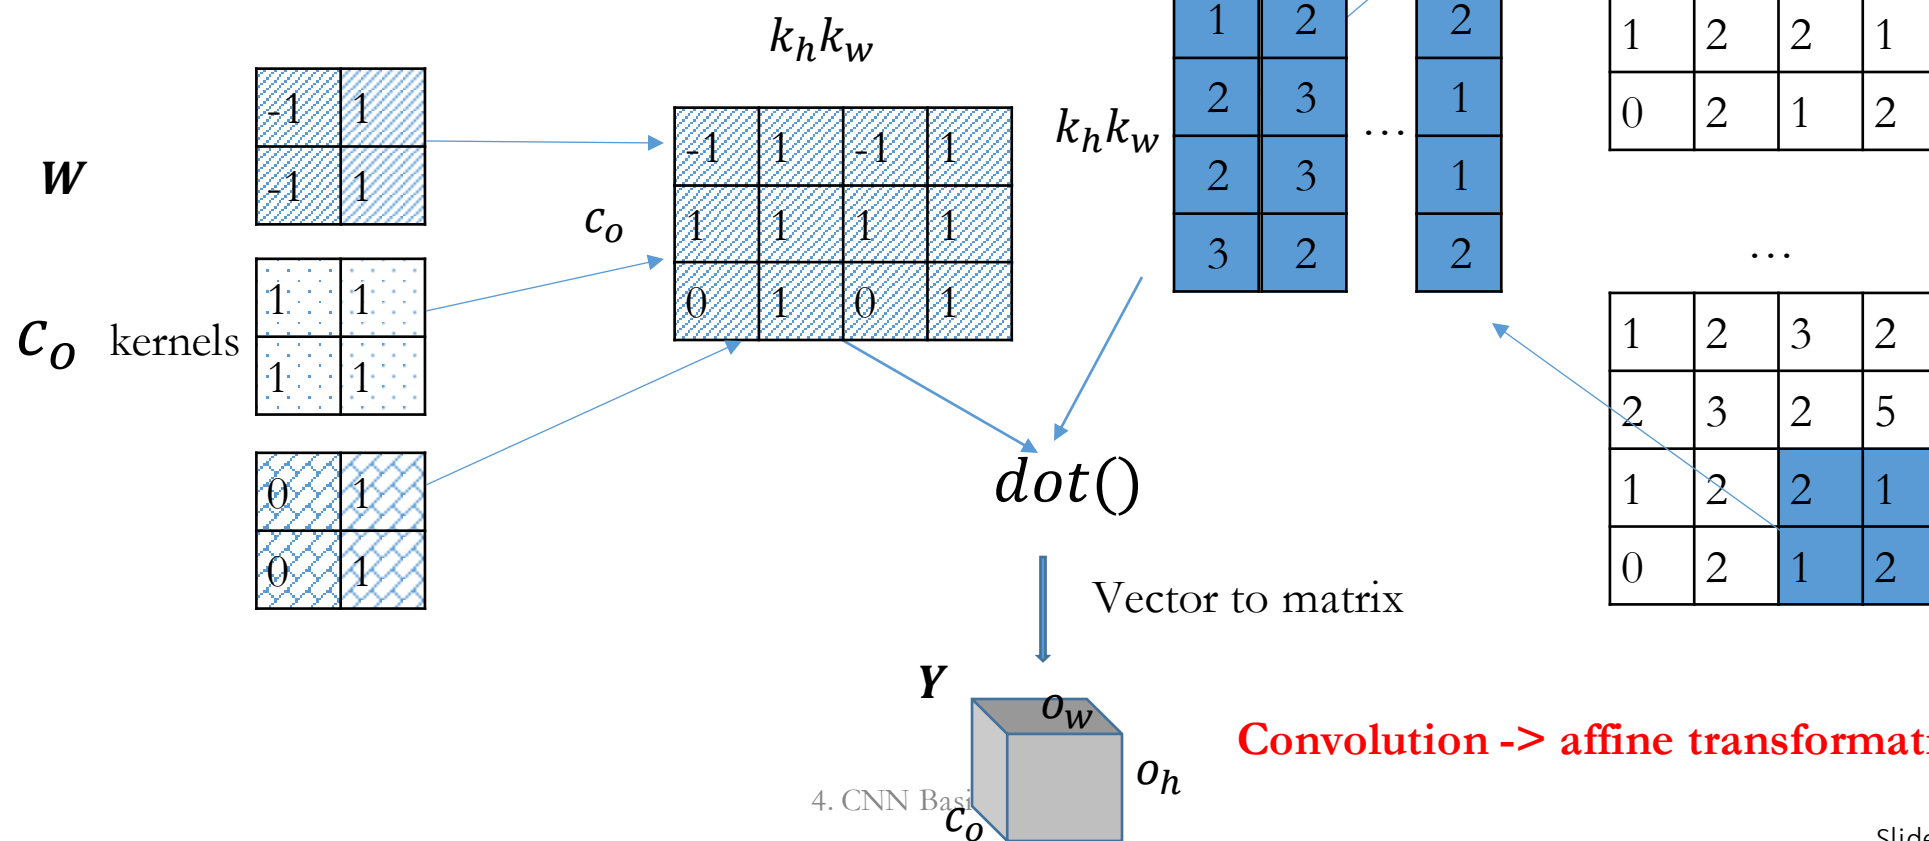
$c_o$  kernels

*Please note,  $c_o$  is not  $C_0$*

$$Y_{l,i,j} = \sum_{a=0}^{k_h-1} \sum_{b=0}^{k_w-1} X_{i*s_h+a,j*s_w+b} \times W_{l,a,b}, l \in [0, c_o)$$

# Implementation

$$Y_{l,i,j} = \sum_{a=0}^{k_h-1} \sum_{b=0}^{k_w-1} X_{i*s_h+a, j*s_w+b} \times W_{l,a,b}, l \in [0, c_o)$$




# Statistics

- applying multiple kernels (filters)  $c_o$ , all of the same stride and padding
- Parameter size
  - $c_o \times k_h \times k_w$
- Output shape
  - $(c_o, o_h, o_w) = (c_o, \left\lfloor \frac{n_h + p_h - k_h}{s_h} \right\rfloor + 1, \left\lfloor \frac{n_w + p_w - k_w}{s_w} \right\rfloor + 1)$
- Computation cost
  - $O((c_o \times k_h \times k_w) \times (o_h \times o_w))$  (float multiplication ops, FLOP)

# 2D Convolution

With multiple ( $c_i$ ) input channels and kernels (filters)

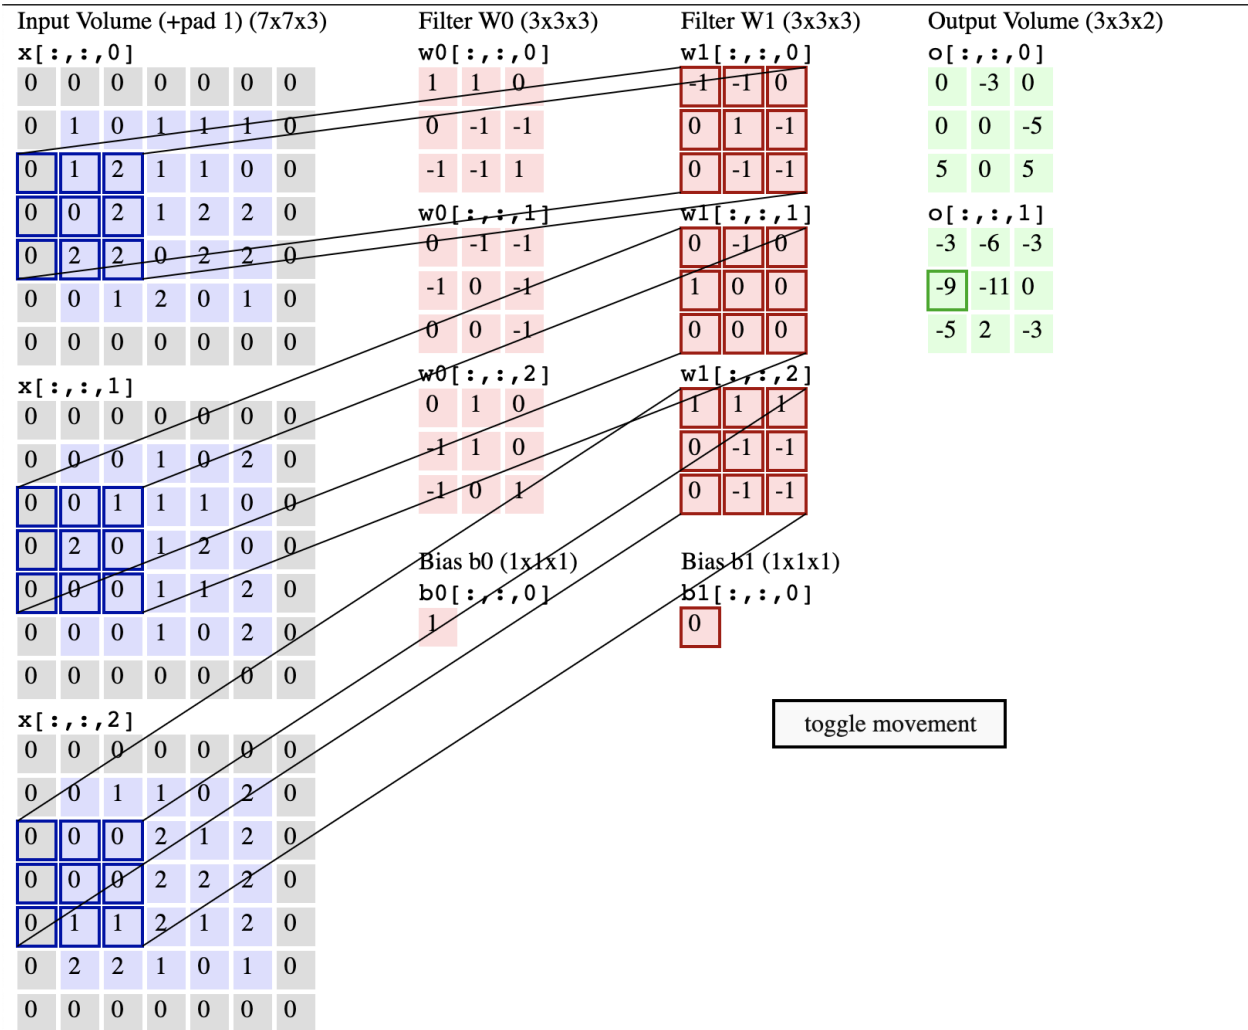


**\*\* The convolution results across all input channels are summed.**

$$Y_{l,i,j} = \sum_{d=0}^{c_i-1} \underbrace{\sum_{a=0}^{k_h-1} \sum_{b=0}^{k_w-1} X_{d,i+a,j+b}}_{\text{convolution result}} \times W_{l,d,a,b} + b_l, l \in [0, c_o) \quad \mathbf{b} \text{ is a bias vector}$$

**\*\* this is still a 2D convolution because the kernel is moved only across the horizontal and vertical dimensions (as indexed by a, b).**

# 2D Convolution (w/ 3D kernels and data)



Demo [\[link\]](#)  
Visualization [\[link\]](#)

toggle movement



# Implementation

- Forward

- Convert input feature maps  $\mathbf{X}$  into matrix  $\hat{\mathbf{X}}$  (img2col) of size  $(c_i k_w k_h \times o_h o_w)$
- Reshape the filters  $\mathbf{W}$  to  $(c_o \times c_i k_h k_w)$
- $\mathbf{Y} = \mathbf{W}\hat{\mathbf{X}} + \mathbf{b}$
- Computational cost,  $O(c_o \times c_i \times k_w \times k_h \times o_h \times o_w)$

- Backward

- Given  $\frac{\partial L}{\partial \mathbf{Y}}$
- Compute  $\frac{\partial L}{\partial \mathbf{W}} = \frac{\partial L}{\partial \mathbf{Y}} \hat{\mathbf{X}}^T$ ,  $\frac{\partial L}{\partial \hat{\mathbf{X}}} = \mathbf{W}^T \frac{\partial L}{\partial \mathbf{Y}}$
- Column to receptive field transformation to get gradient wrt original  $\mathbf{X}$

# Img2Col

- Slide the window from left to right, top to bottom
- Copy the values from the receptive field into a column of  $\hat{\mathbf{X}}$ 
  - Receptive fields across feature maps are concatenated into to one column

- Reshape  $\hat{\mathbf{X}}$  into  $(c_i k_h k_w \times o_h o_w)$

Feature  
map 0

1	2	1	2	3	4
2	2	3	1	2	0
1	1	2	1	0	1
2	1	2	1	1	3

Feature  
map 1

0	2	3	2	0	1
1	0	1	1	2	0
2	0	1	1	1	1
1	1	2	1	1	2




$$c_i k_h k_w$$

$$= 2 \cdot 2 \cdot 2$$

$$o_h o_w = 2 \cdot 3$$

1	1	3	1	2	0
2	2	4	1	1	1
2	3	2	2	2	1
2	1	0	1	1	3
0	3	0	2	1	1
2	2	1	0	1	1
1	1	2	1	2	1
0	1	0	1	1	2

# 2D convolution




# Statistics

- Parameter size
  - Weights:  $c_o \times (c_i k_h k_w)$
  - Bias:  $c_o$
- Output shape
  - $(c_o, o_h, o_w) = (c_o, \left\lfloor \frac{n_h + p_h - k_h}{s_h} \right\rfloor + 1, \left\lfloor \frac{n_w + p_w - k_w}{s_w} \right\rfloor + 1)$
- Computation cost
  - $O(c_o \times c_i k_h k_w \times o_h o_w)$


# Pooling

- Aggregate information from each receptive field
  - Max
  - Average
- No parameters
- Applied for each channel respectively
  - #input channels = # output channels, i.e.,  $c_i = c_o = c$
- Padding and stride can be applied

# Pooling Visualization



# Max Pooling



$C$  feature maps (take max of each one)

1	2	1
<b>3</b>	1	1
-1	0	0

1	1	1
<b>2</b>	0	1
1	0	0

...


0	-2	1
0	1	<b>1</b>
0	0	0

3

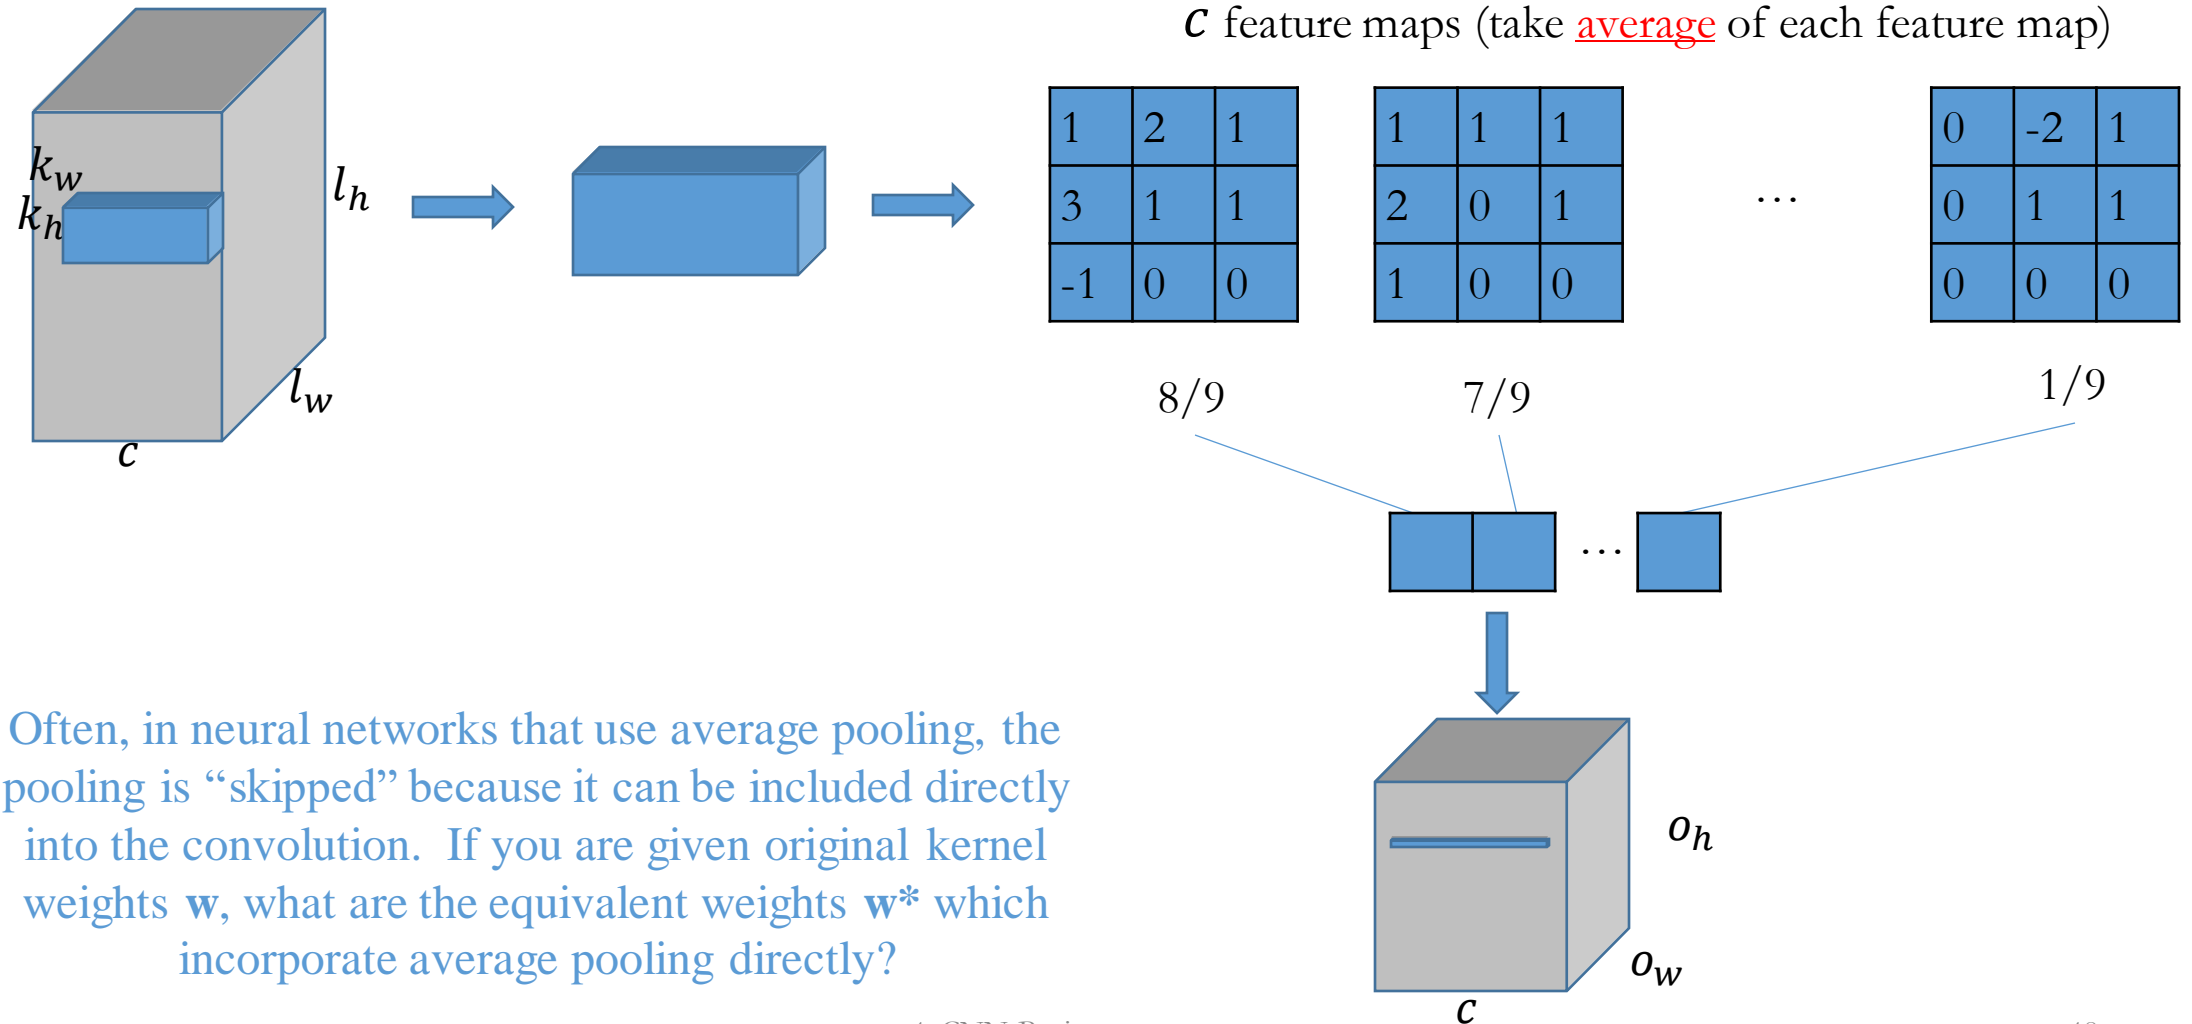
2

1

3	2	...	1
---	---	-----	---



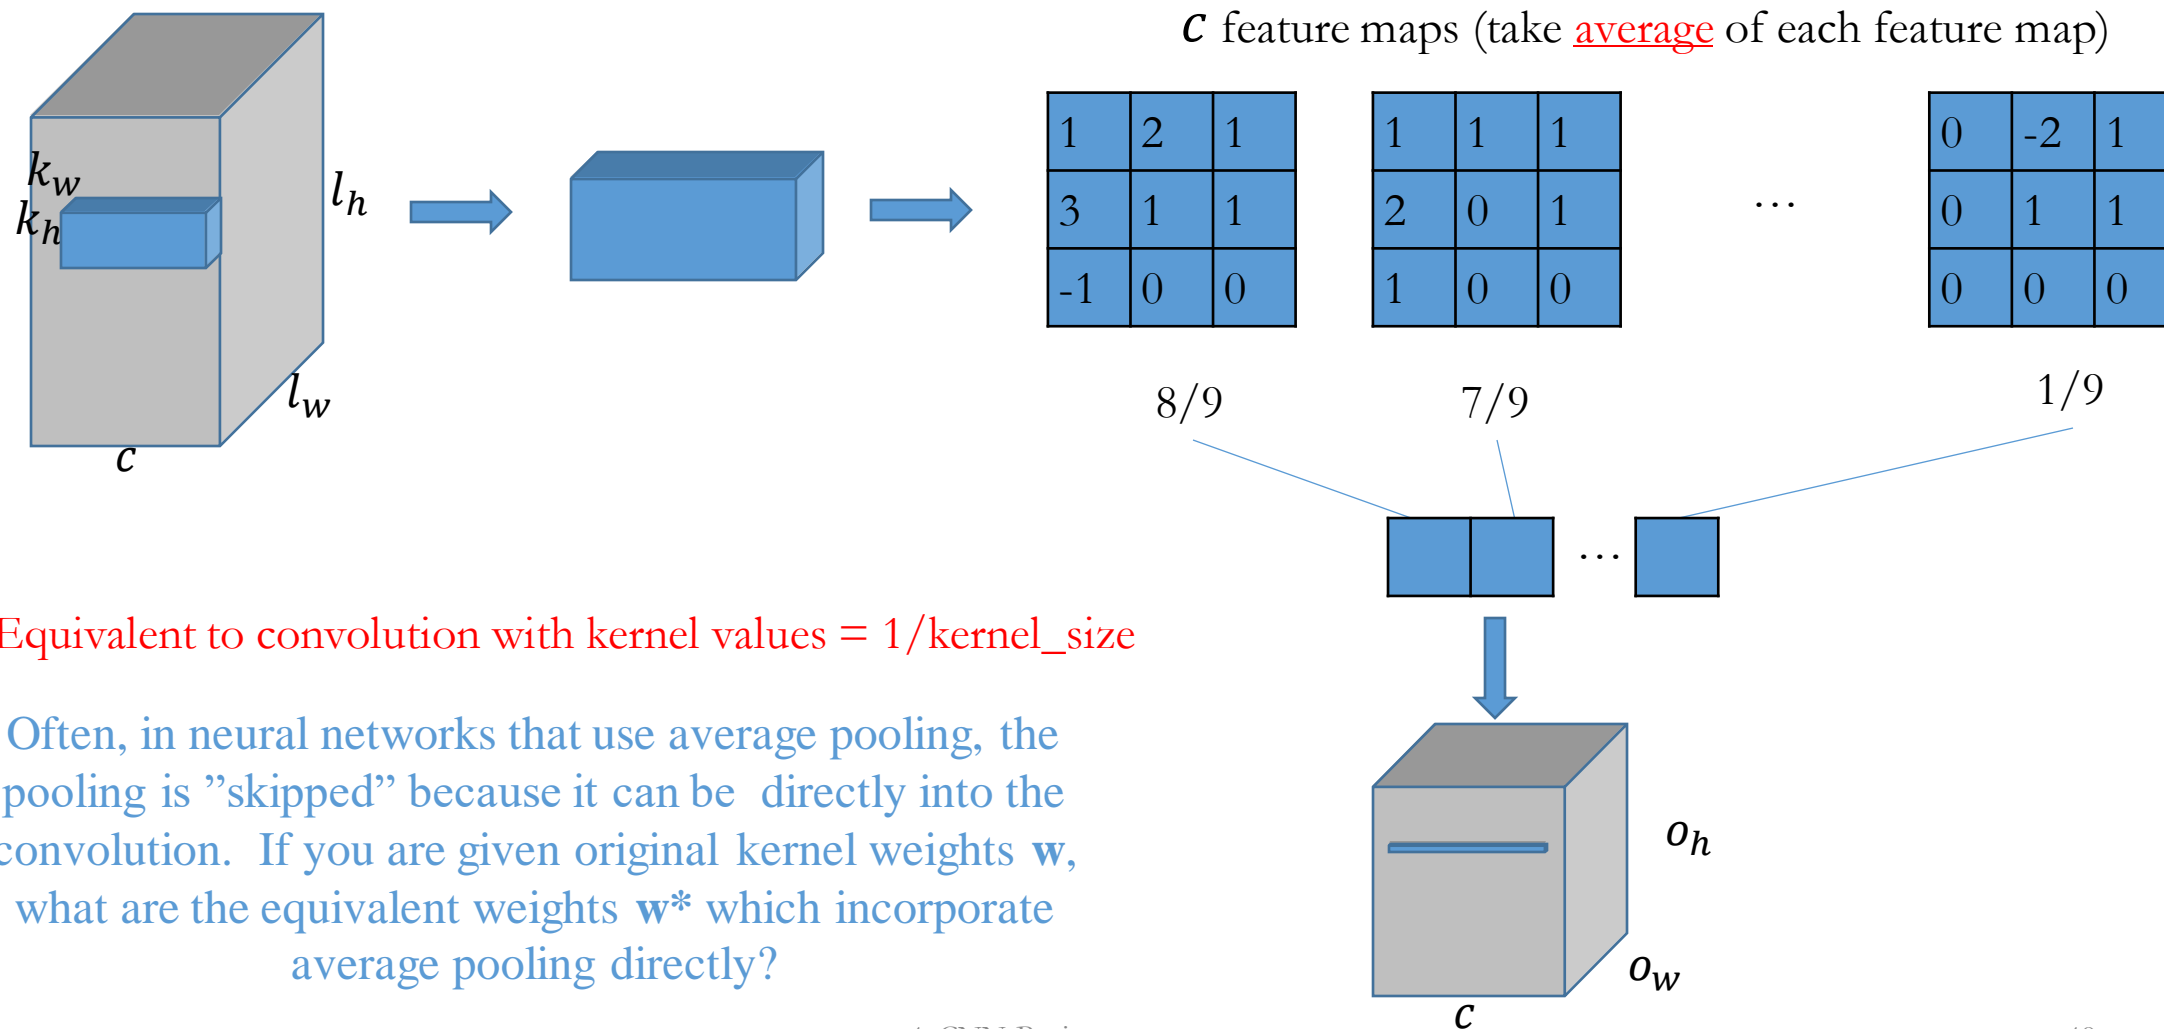
# Average Pooling



Often, in neural networks that use average pooling, the pooling is “skipped” because it can be included directly into the convolution. If you are given original kernel weights  $\mathbf{w}$ , what are the equivalent weights  $\mathbf{w}^*$  which incorporate average pooling directly?




## Average Pooling



# Effect of Pooling

- Reduces the feature size and model size
- Information aggregation
  - Max pooling: invariant to rotation of the **input** image
  - Average pooling: can be replaced by convolution; much cheaper (no weights)



# Multiple input channels, multiple filters

Configuration?

2-minute quiz:


How many input channels?

How many kernels are there?

What is the kernel size?

What is the padding?

What is the stride?



# Multiple input channels, multiple filters

Configuration?


Kernel size:  $3 \times 4 \times 4$

3 input channels

2 filters/kernels

Padding = 3

Stride = 1



# Summary

- Convolution
  - Cross-correlation
  - Kernel, receptive field, padding, stride
  - VS MLP
- 2D convolution
  - Single channel, single kernel
  - Single channel, multiple kernels
  - Multiple channels, multiple kernels
- Pooling
  - Max and Average pooling