

Partially Observable Markov Decision Process

CS4246/CS5446

AI Planning and Decision Making



Please join:
pollev.com/anarayan

This lecture will be
recorded!

About me


Name: *Tasbolat Taunyazov*, P.h.D student
website: www.tasbolat.com



Constrained Orientation Control of a Spherical Parallel Manipulator Via Online Convex Optimization


Tasbolat Taunyazov, Matteo Rubagotti and Almas Shintemirov

ALARIS Astana Laboratory for Robotics and Intelligent Systems
www.alaris.kz



Control

Manipulation



Perception



Topics

- Partially Observable Markov Decision Process (16.4)
 - Belief states and definitions
- Solution algorithms (16.5)
 - Value iteration
 - Online methods (16.5.2 and 3rd ed. 17.4.3)

\

Solving Sequential Decision Problems

- Decision (Planning) Problem or Model

- Appropriate abstraction of states, actions, uncertain effects, goals (wrt costs and values or preferences), and time horizon + observations (through sensing)

- Decision Algorithm

- Input: a problem
- Output: a solution as an optimal action sequence or policy over time horizon

- Decision Solution

- An action sequence or solution from an initial state to the goal state(s)
 - An optional solution or action sequence; OR
 - An optimal policy that specifies “best” action in each state wrt to costs or values or preferences
- (Optional) A goal state that satisfies certain properties

Recall: Decision Making under Uncertainty

- Decision Model:

- **Actions**: $a \in A$
- **Uncertain current state**: $s \in S$ with probability of reaching: $P(s)$
- **Transition model** of uncertain action outcome or effects:
 $P(s' | s, a)$ – probability that action a in state s reaches state s'
- **Outcome** of applying action a :
 $\text{Result}(a)$ – random variable whose values are outcome states
- **Probability of outcome state** s' , conditioning on that action a is executed:
 $P(\text{Result}(a) = s') = \sum_s P(s)P(s' | s, a)$
- **Preferences** captured by a **utility function**:
 $U(s)$ – assigns a single number to express the desirability of a state s

Sequential Decision Problems

- What are sequential decision problems?
 - An agent's utility depends on a sequence of decisions
 - Incorporate utilities, uncertainty, and sensing
 - Search and planning problems are special cases
 - Decision (Planning) Models:
 - Markov decision process (MDP)
 - Partially observable Markov decision process (POMDP)
 - Reinforcement learning: sequential decision making + learning

Quiz

Quiz answer

Recap: Markov Decision Process (MDP)

- Formally:

- An MDP $M \triangleq (S, A, T, R, \gamma)$ consists of:
- A set S of states
- A set A of actions
- A transition function $T: S \times A \times S \rightarrow [0,1]$ that satisfies **Markov Property** such that:

$$\forall s \in S, \forall a \in A: \sum_{\{s' \in S\}} T(s, a, s') = \sum_{\{s' \in S\}} P(s'|s, a) = 1$$

- A reward function $R: S \rightarrow \mathbb{R}$ or $R: S \times A \times S \rightarrow \mathbb{R}$
- A discount factor $0 < \gamma < 1$
- Solution is a **policy** – a function to recommend an action in each state: $\pi: S \rightarrow A$
 - Solution involves careful balancing of risk and reward

Is MDP enough?



Robotics: grasping an item in a clutter
evidence by pixels



Navigation: driving from A to B
evidence by road signs/buildings/etc



Medicine: detecting a diagnosis of a patient
evidence by symptoms

Partially Observable Markov Decision Process (POMDP)

- An POMDP $M \triangleq (S, A, E, T, O, R)$ consists of
- A set S of states
- A set A of actions
- A set E of **evidences** or **observations** or **percepts**
- A transition function $T: S \times A \times S \rightarrow [0,1]$ such that:

$$\forall s \in S, \forall a \in A: \sum_{s' \in S} T(s, a, s') = \sum_{s' \in S} P(s'|s, a) = 1$$

- An **observation function** $O: S \times E \rightarrow [0,1]$ such that:

$$\forall s \in S, \forall e \in E: \sum_{e \in E} O(s, e) = \sum_{e \in E} P(e|s) = 1$$

- A reward function $R: S \rightarrow \mathfrak{R}$ or $R: S \times A \times S \rightarrow \mathfrak{R}$
- Solution: **What is a policy in POMDP?**

Why Study POMDPs?

- **Uncertainty in action outcomes**
 - MDP: fully observable environment
 - Agent knows exactly which state it is in
- **Uncertainty in observations**
 - POMDP: partially observable environment
 - Agent does not know exactly which state it is in – cannot observe the state directly
 - Some noisy observations projected from a state
- **Real-world challenges**
 - Model sequential decision problem for an uncertain, dynamic, partially observable environment
 - If the state is not directly observable, how does an agent reason about its decisions?

Observation Function and Sensor Model

- Observation function:

- $O(s, e) = P(e|s)$ is the probability of observing (or perceiving evidence) e from state s
- Define $O(s, e)$ for all $s \in S$ and $e \in E$
- Assumption of Markov property

- Sensor model

- The observation function defines the sensor model
- An **observation** is also called a **measurement** or **test**

Quiz

Quiz answer

Belief State as Probability Distribution

- Belief State:

- Actual state of the system is unknown, but we can track the probability distribution or **belief state** over the possible states
- An action a changes the belief state b , not just the physical state s
- $b(s)$ denotes probability assigned to actual state s by belief state b
- If $b(s)$ is the current belief, the agent executes action a and receives evidence e' , then the updated belief is given by **filtering**:

$$b'(s') = \alpha \overset{\text{Sensor model}}{P(e'|s')} \sum_s \overset{\text{Transition model}}{P(s'|s, a)} b(s)$$

where α is the normalizing constant that makes the belief state sum to 1

- Filtering Function:

$$b' = \text{FORWARD}(b, a, e')$$

Intuition on Belief State Update

- Recall that Bayes Theorem:

$$P(A|B) = \alpha P(B|A)P(A)$$

where α is normalizing constant.

- Then:

$$\begin{aligned} b'(s') &= P(s'|e', a) && \text{(by definition)} \\ &= \alpha P(e'|s')P(s', a) && \text{(Bayes Theorem)} \\ &= \alpha P(e'|s') \sum_s P(s'|s, a)b(s). && \text{(marginalization over } s) \end{aligned}$$

Exercise

$$b'(s') = \alpha P(e'|s') \sum_s P(s'|s, a) b(s)$$

- Consider a problem with two states s_1 and s_2 :
 - Current belief $b(s_1) = 0.6$ and $b(s_2) = 0.4$.
- For action a :
 - Let the transition probabilities be: $P(s_1|s_1, a) = 0.2$ and $P(s_2|s_1, a) = 0.8$; $P(s_2|s_2, a) = 0.3$ and $P(s_1|s_2, a) = 0.7$
 - Let the observation probabilities be: $P(o_2|s_1) = 0.7$ and $P(o_2|s_2) = 0.2$
- Assume that evidence $e' = o_2$ is received. What is b' ?

$$\begin{aligned} b'(s_1) &= \alpha P(o_2|s_1) [P(s_1|s_1, a) b(s_1) + P(s_1|s_2, a) b(s_2)] \\ &= \alpha 0.7 [0.2 \times 0.6 + 0.7 \times 0.4] = \quad ? \end{aligned}$$

$$\begin{aligned} b'(s_2) &= \alpha P(o_2|s_2) [P(s_2|s_1, a) b(s_1) + P(s_2|s_2, a) b(s_2)] \\ &= \alpha 0.2 [0.8 \times 0.6 + 0.3 \times 0.4] = \quad ? \end{aligned}$$

Decision Making in POMDPs

- Main ideas:

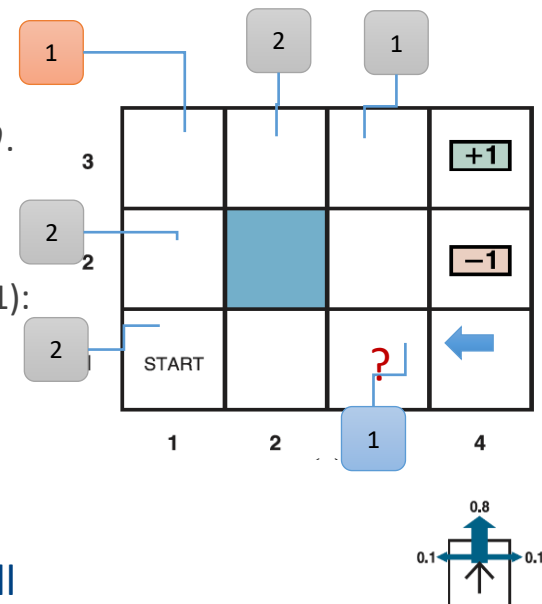
- Optimal action depends only on agent's current belief
- Optimal policy can be described by a mapping $\pi^*(b)$ from belief to action
- Optimal policy does not depend on the actual state agent is in!
- Include **value of information** as a component in decision making

- Decision cycle of a POMDP agent:

- Given current belief b , execute action $a = \pi^*(b)$
- Receive percept e'
- Set belief to $b' = \text{FORWARD}(b, a, e')$ and repeat

Example: Navigation in Grid World

- POMDP :
 - States S , actions A , transition T , reward R , and observation model O .
- Assume:
 - An observation function measures **no. of adjacent walls** in a state s
 - For all non-terminal states except those in column 3 (where value = 1):
 - $O(s, 2) = 0.9, O(s, *) = 0.1$ (where $*$ indicates a wrong value)
- Belief state:
 - $b = \langle 1/9, 1/9, 1/9, 1/9, 1/9, 1/9, 1/9, 1/9, 1/9, 0, 0 \rangle$
 - $b((1, 1)) = 1/9, b((1, 2)) = 1/9, \dots, b((3, 4)) = 0$
- Suppose agent moves Left and sensor reports 1 adjacent wall
 - It is quite likely that agent is now in $(3, 1)$
- What is the exact probability values of the new belief state?



Source: RN Chapter 16

Calculating New Belief State

- What is the probability that an agent in belief state b reaches belief state b' after executing action a ?
 - With current belief $b(s)$, agent executes action a and perceives e' , updated belief is:
$$b'(s') = \alpha P(e'|s') \sum_s P(s'|s, a) b(s) = \text{FORWARD}(b, a, e')$$
where α is normalizing constant for belief state to sum to 1
 - New belief b' is a conditional probability over actual state given sequence of percepts and actions so far
 - If action and subsequent percept were known, deterministic update to the belief using $b' = \text{FORWARD}(b, a, e')$
 - But subsequent percept is not yet known, so agent might arrive in one of several possible belief states b' , depending on percept received

Calculating Probability of Percept

- Probability of percept

- Probability of perceiving e' , given that a was performed starting in belief state b , is given by summing over all the actual states s' that the agent might reach:

$$\begin{aligned} P(e'|a, b) &= \sum_{s'} P(e'|a, s', b) P(s'|a, b) \\ &= \sum_{s'} P(e'|s') P(s'|a, b) \\ &= \sum_{s'} \boxed{P(e'|s')} \sum_s \boxed{P(s'|s, a)} b(s) \end{aligned}$$

Sensor model Transition model

Calculating Transition Model and Reward Model

- Probability of reaching b' from b , given action a :

- Transition model for belief state space:

$$\begin{aligned} P(b'|a, b) &= \sum_{e'} P(b'|e', a, b) P(e'|a, b) \\ &= \sum_{e'} P(b'|e', a, b) \sum_{s'} \underbrace{P(e'|s')}_{\text{Sensor model}} \sum_s \underbrace{P(s'|s, a)}_{\text{Transition model}} b(s). \end{aligned}$$

where $P(b'|e', a, b)$ is 1 if $b' = \text{FORWARD}(b, a, e')$ and 0 otherwise

- Reward function of belief state space:

- Expected reward if the agent does a in belief state b :

$$\rho(b, a) = \sum_s b(s) \sum_{s'} P(s'|s, a) R(s, a, s')$$

Reducing POMDP into an MDP

- Belief space MDP:

- Transition model $P(b'|b, a)$ and reward model $\rho(b, a)$ define an observable MDP on the space of belief states!
- Solving a POMDP on a physical state space – solving a continuous, usually high-dimensional MDP on the corresponding belief-state space
- An optimal policy for this MDP, $\pi^*(b)$ is also an optimal policy for the original POMDP

- Remember:

- The belief state is always observable to the agent, by definition



Value Iteration

Value Iteration for POMPDs

- Policy and conditional plan


- Consider an optimal policy π^* and its application in a specific belief state b
- Policy generates an action, then for each subsequent observation or percept, belief state is updated and new action is generated, and so on
- For b , the policy is equivalent to a conditional plan

- Main issue:

- How does the expected utility of executing a fixed conditional plan varies with the initial belief state?

Conditional Plan

- A policy at a belief b_0 is a conditional plan




- Multiple conditional plans are possible

Utility Function for Belief State


- Note:
 - A belief state b is a probability distribution
 - Each utility value in a POMDP is a function of an entire probability distribution
- Problems:
 - Probability distributions are continuous
 - Huge complexity of belief spaces
- Solution:
 - For finite state, action, and observation spaces and planning horizon: utility functions represented by piecewise linear functions over belief space

Utility Function of Belief State



- Let the utility of executing a fixed conditional plan p starting in physical state s be $\alpha_p(s)$ – **alpha vector**
- Expected utility of executing p in belief state b is:
$$U_p(b) = \sum_s b(s)\alpha_p(s) \text{ or } b \cdot \alpha_p \text{ (inner product of vectors } b, \alpha_p \text{)}$$
 - A linear function of b , corresponding to a hyperplane in belief space

Utility Function of Belief State



- At any particular belief state b , optimal policy is to choose the conditional plan with highest utility:

$$U^{\pi^*}(b) = U(b) = \max_p b \cdot \alpha_p$$

- Utility or value function $U(b)$ on belief states, being the maximum of a collection of hyperplanes, will be **piecewise linear** and **convex**
- For finite depth, there are only a finite number of conditional plans
 - With $|A|$ actions, $|E|$ observations, there are $|A|^{O(|E|^{d-1})}$ distinct depth- d plans

Example: A Two-State World

- Given:

- Two states: A, B (or 0, 1 in RN 3e) [Belief space is 1-dimensional]
- Rewards: $R(.,., A) = 0, R(.,., B) = 1$ [Any transition ending in A is 0, ...]
- Two actions:
 - Stay – Stays put with probability = 0.9
 - Go – Switches to the other state with probability = 0.9
- Discount factor: $\gamma = 1$
- Sensor: Reports correct state with probability = 0.6

- Obviously:

- Agent should: Stay when it thinks it is in state B and Go when it thinks it is in state A

- What are the values for 1-step plans (α -vectors)?

Example: Solving a POMDP, $d = 1$

- Consider one-step plans: [Stay] and [Go]
 - Each receives reward for one transition as follows:

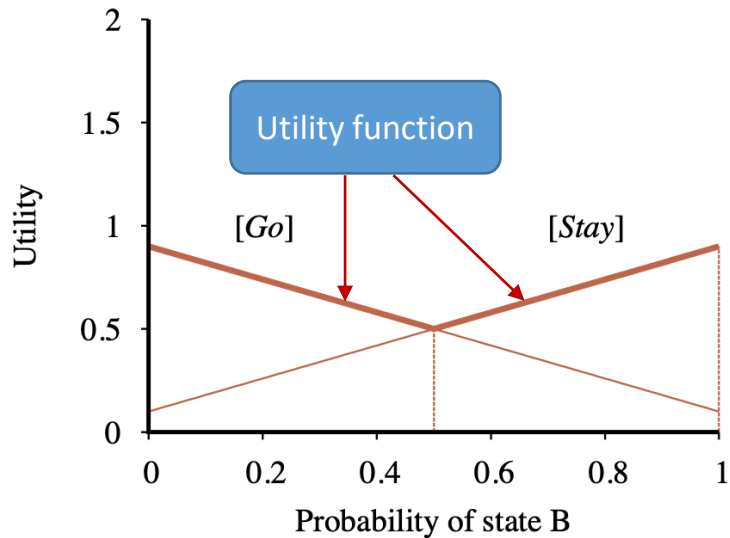
$$\alpha_{[Stay]}(A) = 0.9 R(A, Stay, A) + 0.1 R(A, Stay, B) = 0.1$$

$$\alpha_{[Stay]}(B) = 0.1 R(B, Stay, A) + 0.9 R(B, Stay, B) = 0.9$$

$$\alpha_{[Go]}(A) = 0.1 R(A, Go, A) + 0.9 R(A, Go, B) = 0.9$$

$$\alpha_{[Go]}(B) = 0.9 R(B, Go, A) + 0.1 R(B, Go, B) = 0.1$$

- Hyperplanes for $b \cdot \alpha_{[Stay]}$ and $b \cdot \alpha_{[Go]}$
 - Utility or value function: max of the two linear functions
- What is the one-step optimal policy?
 - Optimal policy is to Stay if $b(B) > 0.5$ and Go otherwise



Utility of Belief-State $b(B)$: $d = 1$


Source: RN Figure 17.15 (a)

Example: Solving a POMDP, $d = 2$


- Deriving a solution:

- Obtain utilities $\alpha_p(s)$ for all the conditional plans p of depth 1 in each physical state s
- Compute utilities for conditional plans of depth 2 by considering:
 - each possible first action
 - each possible subsequent percept, and
 - each way of choosing a depth-1 plan to execute for each percept:
 - [Stay; **if** Percept = A **then** Stay **else** Stay]
 - [Stay; **if** Percept = A **then** Stay **else** Go]
 - [Go; **if** Percept = A **then** Stay **else** Stay] ...
- There are 8 distinct depth-2 plans in all
 - 4 suboptimal and **dominated plans** – dominated plans are never optimal
 - 4 **undominated plans**, each optimal in a specific region
- The regions partition the belief-state space

Example: Utility of Belief-State $b(B): d = 2$



Utilities of 8 distinct 2-step plans




Utilities of 4 undominated 2-step plans

- Continuous belief space is divided into regions; the regions partition the belief-state space
 - Each region corresponds to a conditional plan that is optimal for that region
- $U(b)$ is piecewise linear and convex

Source: RN Figure 16.15 (b) and (c)

Value Iteration in POMDP



- In general:

- Let p be a depth- d conditional plan with initial action a followed by depth $(d - 1)$ subplans $p \cdot e'$ for percept e'

$$\alpha_p(s) = \sum_{s'} P(s'|s, a) \left[R(s, a, s') + \gamma \sum_{e'} P(e'|s') \alpha_{p \cdot e'}(s') \right]$$

Transition model
Observation model


- Given utility function:

- Executable policy extracted by:
 - looking at which hyperplane is optimal at any given belief state b
 - executing the first action of the corresponding plan

- POMDP Value Iteration (VI):

- maintains a collection of undominated plans $\{p\}$ with their utility hyperplanes $\{\text{alpha vectors } \alpha_p\}$
- Cf MDP VI computes one utility number, $U(s)$ for each state s

Example: Utility of Belief-State $b(B)$: $d = 8$



Utility function for optimal 8-step plans

Source: RN Figure 17.15 (d)

POMDP Value Iteration Algorithm

function POMDP-VALUE-ITERATION(*pomdp*, ϵ) **returns** a utility function
 inputs: *pomdp*, a POMDP with states S , actions $A(s)$, transition model $P(s' | s, a)$,
 sensor model $P(e | s)$, rewards $R(s, a, s')$, discount γ
 ϵ , the maximum error allowed in the utility of any state
 local variables: U , U' , sets of plans p with associated utility vectors α_p

 $U' \leftarrow$ a set containing all one-step plans $[a]$, with $\alpha_{[a]}(s) = \sum_{s'} P(s' | s, a) R(s, a, s')$
 repeat
 $U \leftarrow U'$
 $U' \leftarrow$ the set of all plans consisting of an action and, for each possible next percept,
 a plan in U with utility vectors computed according to Equation (16.18)
 $U' \leftarrow \text{REMOVE-DOMINATED-PLANS}(U')$
 until MAX-DIFFERENCE(U, U') $\leq \epsilon(1 - \gamma)/\gamma$
 return U

Source: RN Figure 16.16

Notes on Value Iteration

- Algorithm's complexity depends primarily on how many plans get generated.
 - Given $|A|$ actions and $|E|$ observations, there are $|A|^{O(|E|^{d-1})}$ distinct depth- d plans.
 - $|A|n^{|E|}$ conditional plans generated at level $d + 1$ before elimination, where n is the no. of conditional plans at level d
 - Elimination of dominated plans to reduce doubly exponential growth
 - P-SPACE hard – very inefficient in practice
 - Approximate methods necessary
- Intermediate belief states have lower value than state A and state B
 - In the intermediate states the agent lacks information needed to choose a good action.
 - Information has value and optimal policies in POMDPs often include information-gathering actions.



Online Methods

Approximate solutions



Scaling POMDP solvers

- POMDP solvers need to solve two problems
 - Belief tracking/filtering
 - Given the history observed, what is the current belief?
 - Planning
 - Given the current belief, what is the optimal action to take?
- To scale up, need to scale up for both problems

Online Agent for POMDP

- Main ideas:

- Represent transition model and sensor model by a **dynamic decision network (DDN)**
- Belief tracking – Inference in DDN – Exact inference is computationally intractable
 - No known polynomial time algorithm, as number of state variables grow
- Approximate solution: Deploy a **filtering** algorithm (exact or particle filtering) to incorporate each new percept and action and to update belief state representation
- Projecting forward possible action sequences and choosing the best one

- Note:

- A DDN can actually be used as inputs for any POMDP solution algorithm

Dynamic Decision Network (DDN)

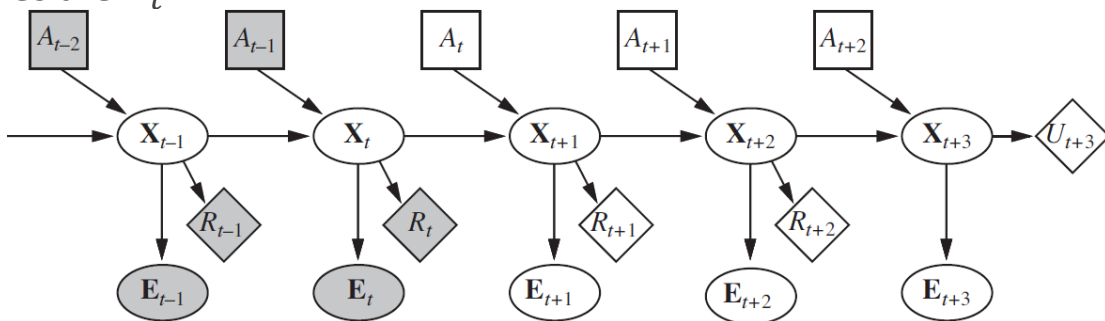
- Execution of POMDP over time can be represented as a DDN
 - Transition and sensor models represented by a Dynamic Bayesian Network (DBN)
 - Add decision and utility nodes to get DDN

- In DDN:



- State S_t becomes set of variables \mathbf{X}_t
- Evidence/observation variables are \mathbf{E}_t
- Action at time t is A_t
- Transition: $P(\mathbf{X}_{t+1}|\mathbf{X}_t, A_t)$
- Sensor model: $P(\mathbf{E}_t|\mathbf{X}_t)$

Note:

- Variables with known values are shaded
- Current time is t and agent must decide what to do



Look-Ahead Solution of POMDP



Alternate form

Time complexity: $O(|A|^d |E|^d)$

Notes on Online Solution

- Main ideas:

- Project action sequences forward from current belief state
- Determine utility estimates at projected depth based on future rewards
- Extract decision by backing up utilities from the leaves
 - Averaging chance nodes and maximizing/minimizing decision nodes
- Non-leaf states may have rewards
- Decision nodes correspond to belief states rather than actual states

- Complexity

- Time complexity of exhaustive search to depth d is $O(|A|^d |E|^d)$
where $|A|$ is the no. of available actions and $|E|$ is the number of possible percepts

POMCP¹

*Partially Observable
Monte Carlo Planning*

MCTS:


- Select
- Expand
- Simulate
- Backup

- To run UCT on POMDP, we need to represent beliefs in the nodes
 - Inference to construct beliefs is intractable in general
 - For MDP, states are easy to generate
 - Beliefs are difficult to generate, so POMCP uses action-observation history
 - History is equivalent to belief assuming same initial belief; $|A|^d k^d$
- Instead of propagating beliefs forward in a trial
 - POMCP samples a state at the root from the initial belief.
 - Run the simulation using the state to generate action-observation history
 - Only need to sample from $P(s'|s, a)$ then $P(e'|s')$ to generate observation e' for the action-observation history
 - Avoid constructing beliefs

Silver, D. and J. Veness, *Monte-Carlo planning in large POMDPs*, in *Advances in Neural Information Processing Systems 23: 24th Annual Conference on Neural Information Processing Systems 2010*. . 2010: Vancouver, British Columbia, Canada. p. 2164-2172

DESPOT²

*Determinized Sparse
Partially Observable Tree*



Source: Ye et al 2017

- Construct search tree differently! (make tree smaller)
 - Construct k different search trees
 - Each search tree, sample a state at the root to initialize
 - At every action node, try all actions on the (single) state at that node
 - At every observation node, sample a single observation from the (single) state at node
- Size of the tree:
 - Each of the k trees have size A^d . Combine together to get tree of size $O(|A|^d k)$
 - Exponentially smaller than $|A|^d |E|^d$
- Can show that searching sampled tree is sufficient if a small good policy exists
 - Use heuristics to do anytime search of this tree

²Ye, N., et al., *DESPOT: online POMDP planning with regularization*. J. Artif. Int. Res., 2017. **58**(1): p. 231–266.

POMDP Applications

- Autonomous driving golf cart in UTown
 - https://youtu.be/y_9VMD_sQhw
 - DESPOT is used there
 - Works in real-time
- Robotics
- Airborne Collision Avoidance System X
 - ACAS X
 - POMDPs with neural networks for function approximation
- Education
- Dialog system/Chatbot
 - Assistive agent for dementia patients



ARTICLE

Automated handwashing assistance for persons with dementia using video and a partially observable Markov decision process

Twitter LinkedIn YouTube Facebook Email

Authors: Jesse Hoey, Pascal Poupart, Axel von Bertoldi, Tammy Craig, Craig Boutilier, Alex Mihailidis [Authors Info & Claims](#)

Computer Vision and Image Understanding, Volume 114, Issue 5 • May, 2010 • pp 503–519 • <https://doi.org/10.1016/j.cviu.2009.06.008>

Published online 1 May 2010

Summary

- POMDPs

- Compute optimal policy under partial observability.
- For finite horizon problems, resulting optimal utility (value) functions are continuous, piecewise linear, and convex.
- In each iteration, no. of linear constraints grows exponentially.
- Approximate solvers can scale to large size problems; solvers mainly remain in research settings
 - For example, the SARSOP solver from NUS does the Bellman update only on a small subset of beliefs and use heuristics to explore the belief space
 - <https://www.comp.nus.edu.sg/~leews/publications/rss08.pdf>
- Online search tends to scale better.

Homework

- Readings

- RN 16.4, 16.5
- RN 14.2.1, 14.5.3 (*Filtering and Particle Filtering – Optional*)

- References

- Kaelbling, L.P., M.L. Littman, and A.R. Cassandra, *Planning and acting in partially observable stochastic domains*. Artif. Intell., 1998. **101**(1–2): p. 99–134.
- Silver, D. and J. Veness, *Monte-Carlo planning in large POMDPs*, in *Advances in Neural Information Processing Systems 23: 24th Annual Conference on Neural Information Processing Systems 2010*. . 2010: Vancouver, British Columbia, Canada. p. 2164-2172.
- Ye, N., et al., *DESPOT: online POMDP planning with regularization*. J. Artif. Int. Res., 2017. **58**(1): p. 231–266.