

**CS4347**

# **Sound and Music Computing**

**L5: Automatic Speech Recognition (ASR)**

**Wang Ye**

**[www.comp.nus.edu.sg/~wangye](http://www.comp.nus.edu.sg/~wangye)**

**[wangye@comp.nus.edu.sg](mailto:wangye@comp.nus.edu.sg)**

**Office: AS6-04-08**

# Topics to Cover (selective approach)

## Part A: The Core

- Introduction
- Review of DFT, Audio Representation, and Machine Learning
- Music Representation, Analysis and Transcription
- Automatic Music Transcription (AMT)
- Automatic Speech Recognition (ASR)
- Generative Models for Text-to-Speech (TTS) & Singing Voice Synthesis (SVS)

 Midterm break

## Part B: The Breadth

- Spoken language assessment
- Singing voice processing
- Music production audio effects
- Automatic Music Generation
- Synthesis of sound & music – a DSP approach
- Project presentations/demo

# Topics Today

➡ Part A: Overview of Automatic Speech Recognition (ASR)

Part B: Development of ASR System

Part C: End-to-End ASR System

Part D: Automatic Lyric Transcription (ALT)

# What is Automatic Speech Recognition (ASR)?



- ASR task aims to transcribe a speech waveform into a text transcript
- Also known as speech-to-text (STT); the reverse is known as TTS
- Variants – keyword spotting (KWS), voice command recognition
- Speaker identification/verification
- Variants – e.g, speaker diarization
- Large vocabulary continuous speech recognition (LVCSR)


Draw an analogy between speech and **music**

# Speech is usually represented as a sequence of acoustic features




Acoustic Features

$$X = \{x_t \in R^D | t = 1, 2, \dots, T\}$$



Raw waveform:

- 1-D Temporal domain representation
- Sampling rate, channels
- **High-dimensional**



Spectral features:

- Time-frequency domain representation
- MFCC, STFT, filter bank
- **Low-dimensional**

# Text transcript is represented as a sequence of tokens



Text transcript

$$W = \{w_n \in V | n = 1, 2, \dots, N\}$$

IF I DO NOT BELIEVE IN DOGMA IT IS BECAUSE I BELIEVE IN FREEDOM

Vocabulary (types of tokens):

- Word
- Character
- Phoneme
- Sub-word

# Text transcript is represented as a sequence of tokens

*i*

Text transcript

$$W = \{w_n \in V | n = 1, 2, \dots, N\}$$

IF I DO NOT BELIEVE IN DOGMA IT IS BECAUSE I BELIEVE IN FREEDOM



$w_i$



$w_j$

Vocabulary:

- Word
- Character
- Phoneme
- Sub-word

# Text transcript is represented as a sequence of tokens



Text transcript

$$W = \{w_n \in V | n = 1, 2, \dots, N\}$$

IF I DO NOT BELIEVE IN **D**OGMA IT IS **B**ECAUSE I BELIEVE IN FREEDOM

└

$w_i$

└

$w_j$

Vocabulary:

- Word
- **Character**
- Phoneme
- Sub-word



# Text transcript is represented as a sequence of tokens



Text transcript

$$W = \{w_n \in V | n = 1, 2, \dots, N\}$$

IF I DO NOT BELIEVE IN DOGMA IT IS BECAUSE I BELIEVE IN FREEDOM

$/n/$   $/ao/$   $/t/$   
└─┘  
 $w_i$

$/iy/$   $/z/$   
└─┘  
 $w_j$

Vocabulary:

- Word
- Character
- **Phoneme**
- Sub-word


# Text transcript is represented as a sequence of tokens



Text transcript

$$W = \{w_n \in V | n = 1, 2, \dots, N\}$$

IF I DO NOT BELIEVE IN DOGMA IT IS BECAUSE I BELIEVE IN FREEDOM




Vocabulary:

- Word
- Character
- Phoneme
- **Sub-word**

# How to build an ASR system?


Modeling ASR task as a machine learning problem:

- Supervised learning dominates the paradigm



Training phase

Testing phase



- Semi-supervised and unsupervised learning are being investigated for ASR

# How to evaluate an ASR system?

Word Error Rate (WER): most widely used metric

Ref.

IF I DO NOT BELIEVE IN DOGMA IT IS BECAUSE I BELIEVE IN FREEDOM

Out.

IF I DO NOT BELIEVE IN DOGMA IT IS BECAUSE I BELIEVE IN KINGDOM

*S*

Substitutions

# How to evaluate an ASR system?

Word Error Rate (WER): most widely used metric

Ref.

IF I DO NOT BELIEVE IN DOGMA IT **IS** BECAUSE I BELIEVE IN FREEDOM

Out.

IF I DO NOT BELIEVE IN DOGMA IT      BECAUSE I BELIEVE IN FREEDOM

*D*

Deletions

# How to evaluate an ASR system?

Word Error Rate (WER): most widely used metric

Ref. IF I DO NOT BELIEVE IN DOGMA IT IS BECAUSE I BELIEVE IN FREEDOM

Out. IF I DO NOT BELIEVE IN DOGMA IT IS BECAUSE I **AM** BELIEVE IN FREEDOM

*I*

**Insertions**

# How to evaluate an ASR system?

Word Error Rate (WER): most widely used metric

WER ↓ ↔ Performance ↑

$$WER = \frac{S + D + I}{N} = \frac{S + D + I}{S + D + C}$$

- $S$  is the number of word substitutions
- $D$  is the number of word deletions
- $I$  is the number of word insertions
- $C$  is the number of correct words
- $N$  is the total number of words in the reference

Similarly, we can define:

Character Error Rate (CER)

Phoneme Error Rate (PER), etc.

# Topics Today

Part A: Overview of Automatic Speech Recognition (ASR)

 Part B: Development of ASR System


Part C: End-to-End ASR System

Part D: Automatic Lyric Transcription (ALT)



# Development of ASR System over the years

ASR Performance on Switchboard dataset (Benchmark English Corpus)



HMM/GMM framework dominated ASR for half a century

DL based methods have emerged in the past decade

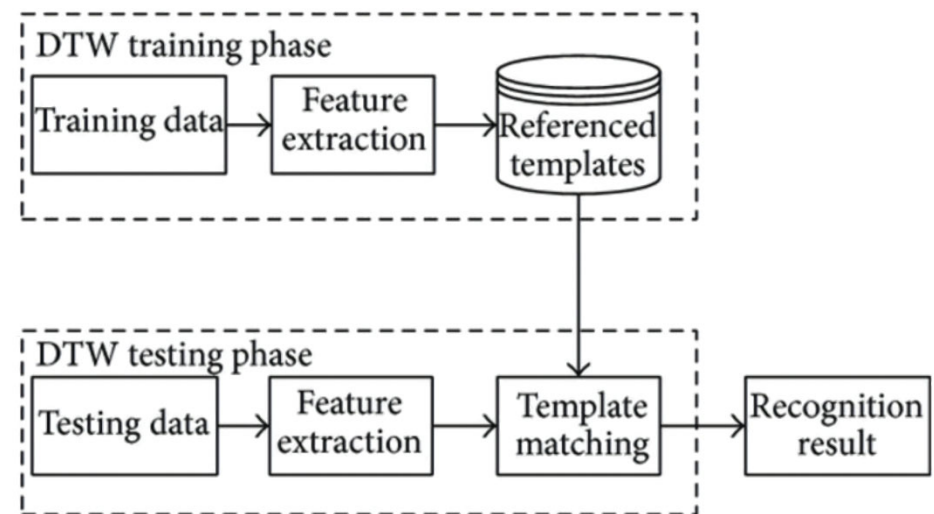
# Template-based ASR

Training phase:

- Build a database for referenced templates
- Each template is the representation of an actual speech segment together with its transcription, neighboring templates, meta-information, etc.

Testing phase:

- Match the speech segments of test data with referenced templates through DTW (Dynamic Time Warping) algorithm




***e.g., digit recognition***

# Statistical ASR

Statistical Approach:


- During the testing (or called decoding): Aims to find the most likely sequence of tokens ( $W$ ) given the sequence of acoustic features ( $X$ )



# Statistical ASR

Statistical Approach:

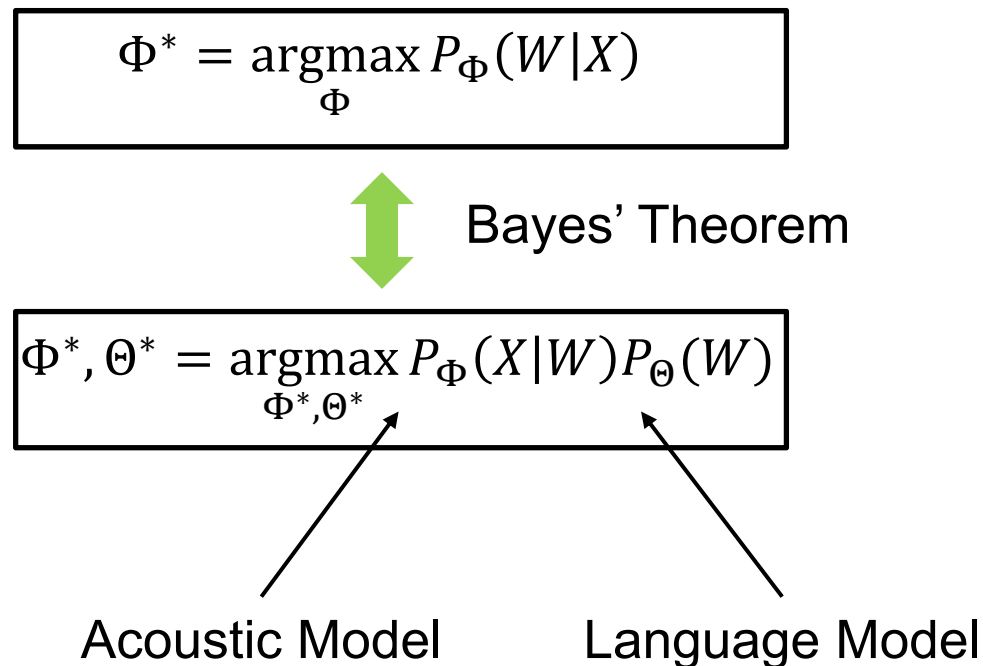
- During the testing (or called decoding): Aims to find the most likely sequence of tokens ( $W$ ) given the sequence of acoustic features ( $X$ )
- During the training: Aims to train the model  $\Phi$  to maximize the probability given the pairs of  $W$  and  $X$



# Statistical ASR: HMM-based ASR

Hidden Markov Model (HMM)-based ASR:

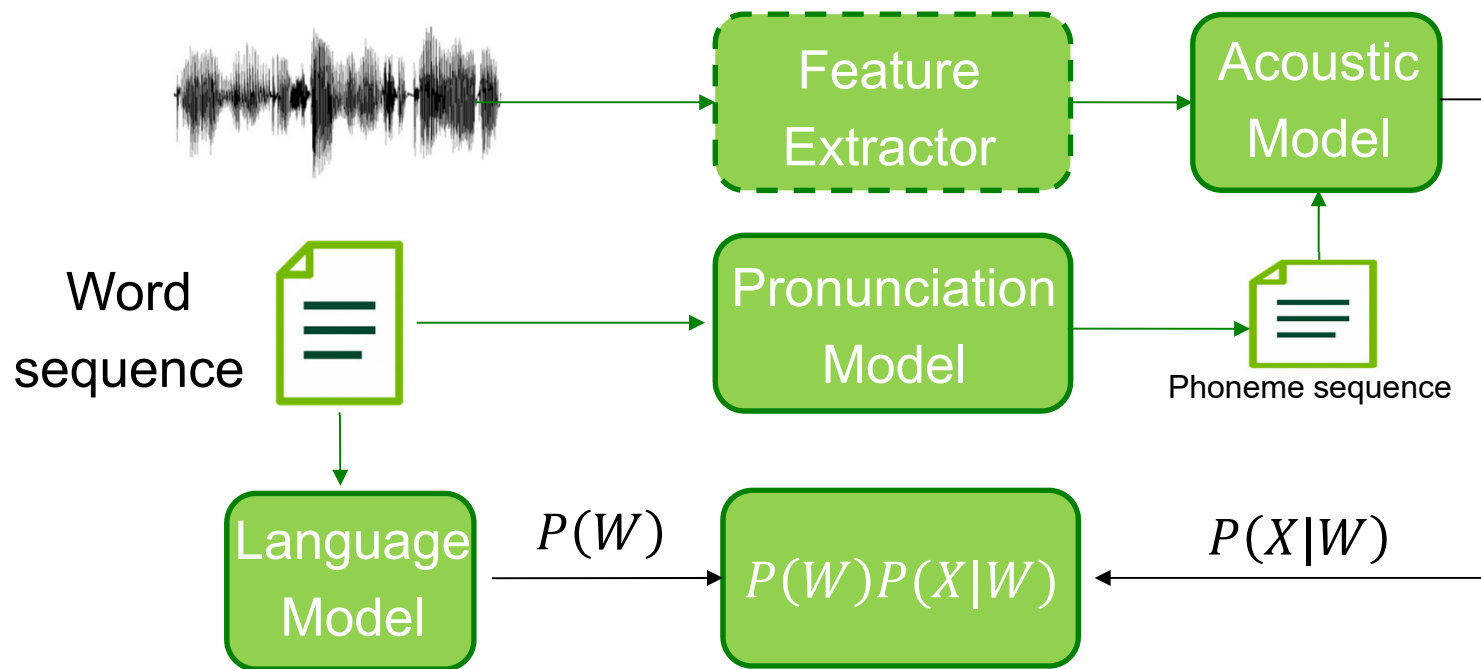
- Adopt Bayes' Theorem to optimize the **posterior**
- Model the **likelihood** via Acoustic Model and **prior** via Language Model




# Statistical ASR: HMM-based ASR

HMM-based ASR:

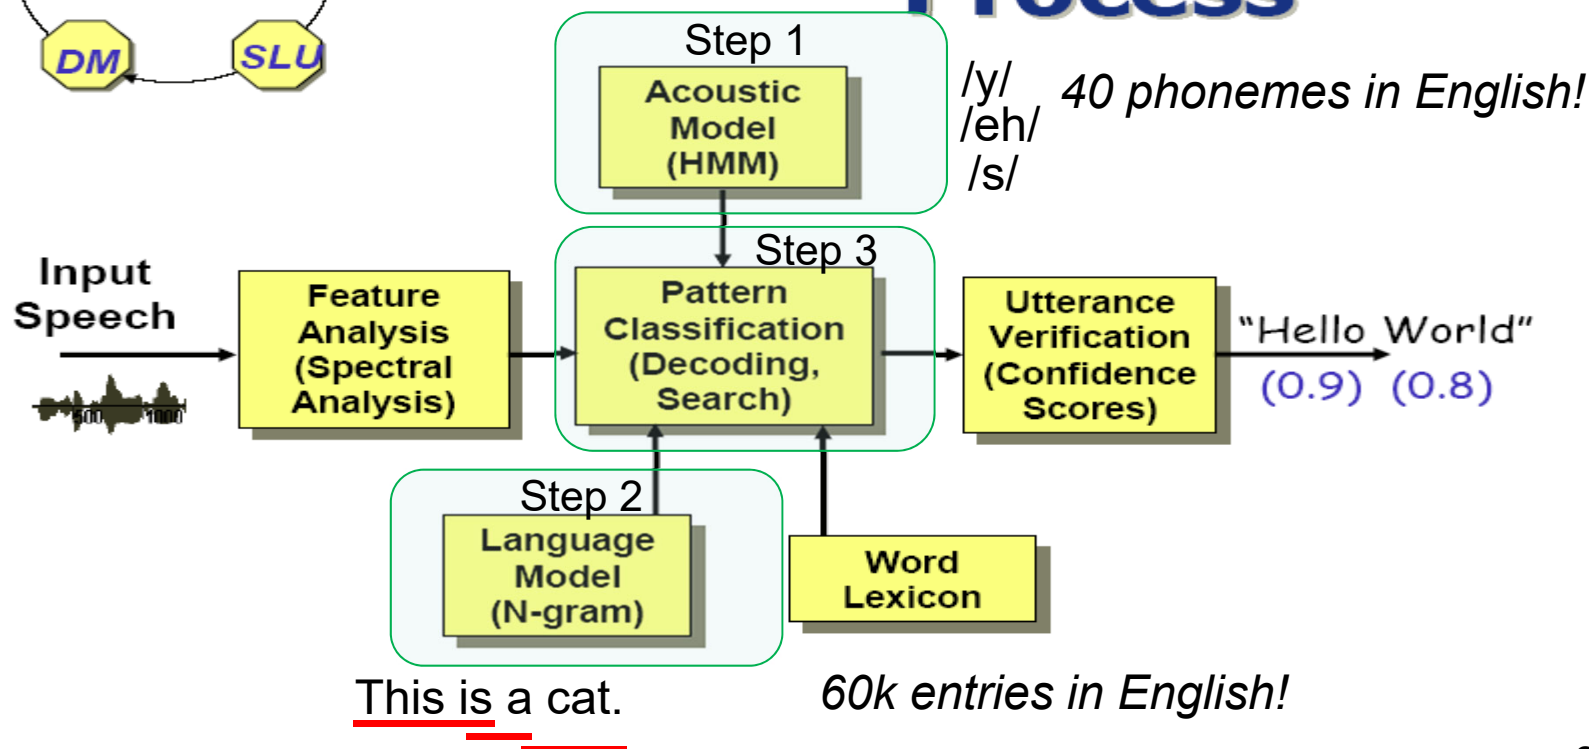
- Adopt Bayes' Theorem to optimize the posterior
- Model the likelihood via Acoustic Model and prior via Language Model
- Split word sequence into phoneme sequence via Pronunciation Model



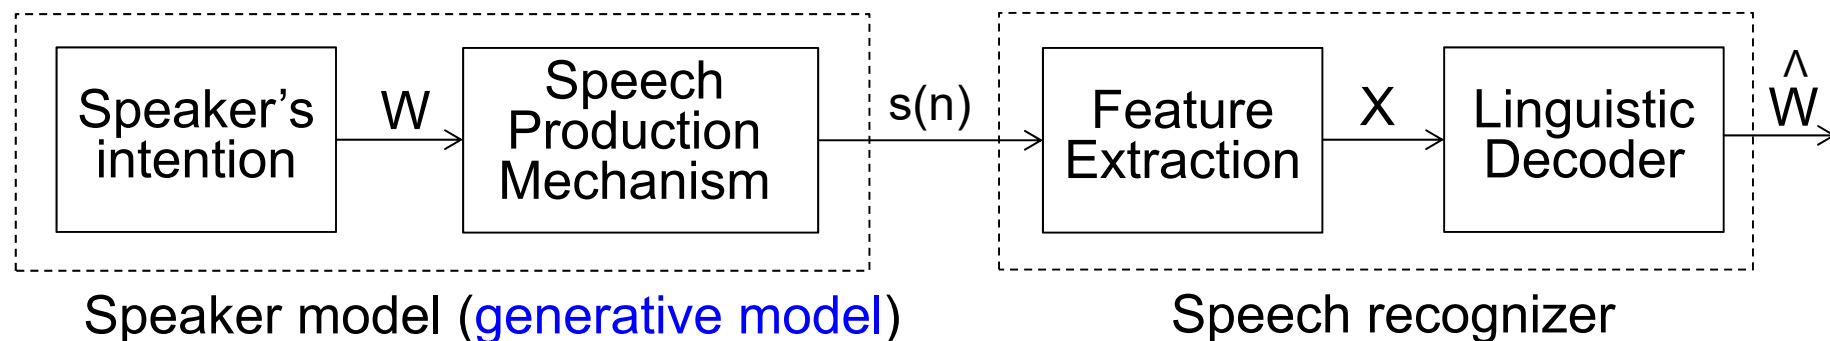
# HMM based Automatic Speech Recognition (ASR) System



## Speech Recognition Process



# Basic ASR Formulation (Bayes Method)



$$\begin{aligned}
 \hat{W} &= \arg \max_W P(W | X) && \text{Posterior probability (X produced by W)} \\
 &= \arg \max_W \frac{P(X | W) P(W)}{P(X)} && \begin{array}{l} \text{Probability that a given model (W) produces X} \\ \text{(acoustic model) – it is also called likelihood} \end{array} \\
 &= \arg \max_W \underbrace{P_A(X | W)}_{\text{Step 1}} \underbrace{P_L(W)}_{\text{Step 2}} && \begin{array}{l} \text{Prior probability of the model (language model)} \\ \text{Prior probability of the observation} \end{array}
 \end{aligned}$$


Step 3 is indicated by a bracket under the  $\arg \max_W$  term.



# Statistical ASR: HMM-based ASR

HMM-based ASR:

- Acoustic Model is HMM-based
- State can be the phoneme  $s_t$

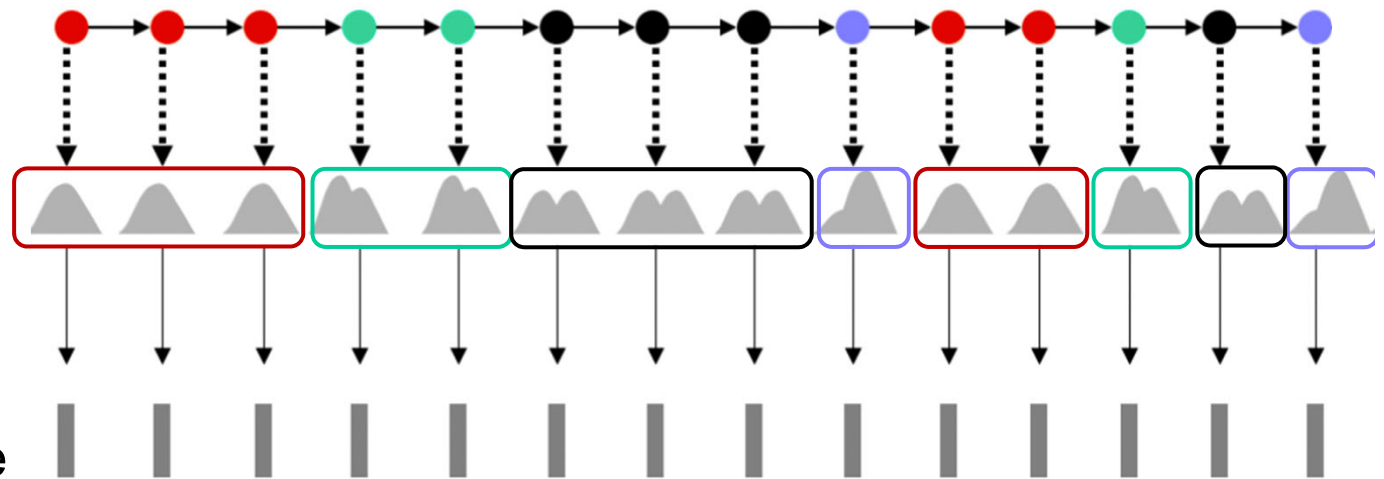


State transition  
 $p(s_t | s_{t-1})$

**State sequence  
(hidden)**  $s_t$

*Emission distribution*  
 $p(x_t | s_t)$

**Observation sequence  
(feature vectors)**  $x_t$



# Statistical ASR: HMM-based ASR

## GMM-HMM ASR

- Model the emission probability density via Gaussian Mixture Model (GMM)
- GMM is a linear combination of Gaussian distribution

$$p(\mathbf{x}) = \sum_{m=1}^M P(m)p(\mathbf{x}|m) = \sum_{m=1}^M P(m)N(\mathbf{x}; \boldsymbol{\mu}_m, \sigma_m^2 \mathbf{I})$$

## DNN-HMM ASR

- Model the emission probability density via Deep Neural Network (DNN)

# Statistical ASR: End-to-End ASR

End-to-End ASR:

- Directly optimize the following objective
- Achieve State-of-the-art Performance on benchmark ASR datasets

The diagram illustrates the end-to-end ASR objective. At the top, a rectangular box contains the equation  $\Phi^* = \underset{\Phi}{\operatorname{argmax}} P_{\Phi}(W|X)$ . Two arrows originate from below the box: one points to the  $W$  term and the other points to the  $X$  term. Below the left arrow, the text 'Word sequence' is followed by the set definition  $W = \{w_n \in V | n = 1, 2, \dots, N\}$ . Below the right arrow, the text 'Acoustic features' is followed by the set definition  $X = \{\mathbf{x}_t \in R^D | t = 1, 2, \dots, T\}$ .

$$\Phi^* = \underset{\Phi}{\operatorname{argmax}} P_{\Phi}(W|X)$$

Word sequence  
 $W = \{w_n \in V | n = 1, 2, \dots, N\}$


Acoustic features  
 $X = \{\mathbf{x}_t \in R^D | t = 1, 2, \dots, T\}$

Focus of this lecture and assignment 3!

# Topics Today

Part A: Overview of Automatic Speech Recognition (ASR)

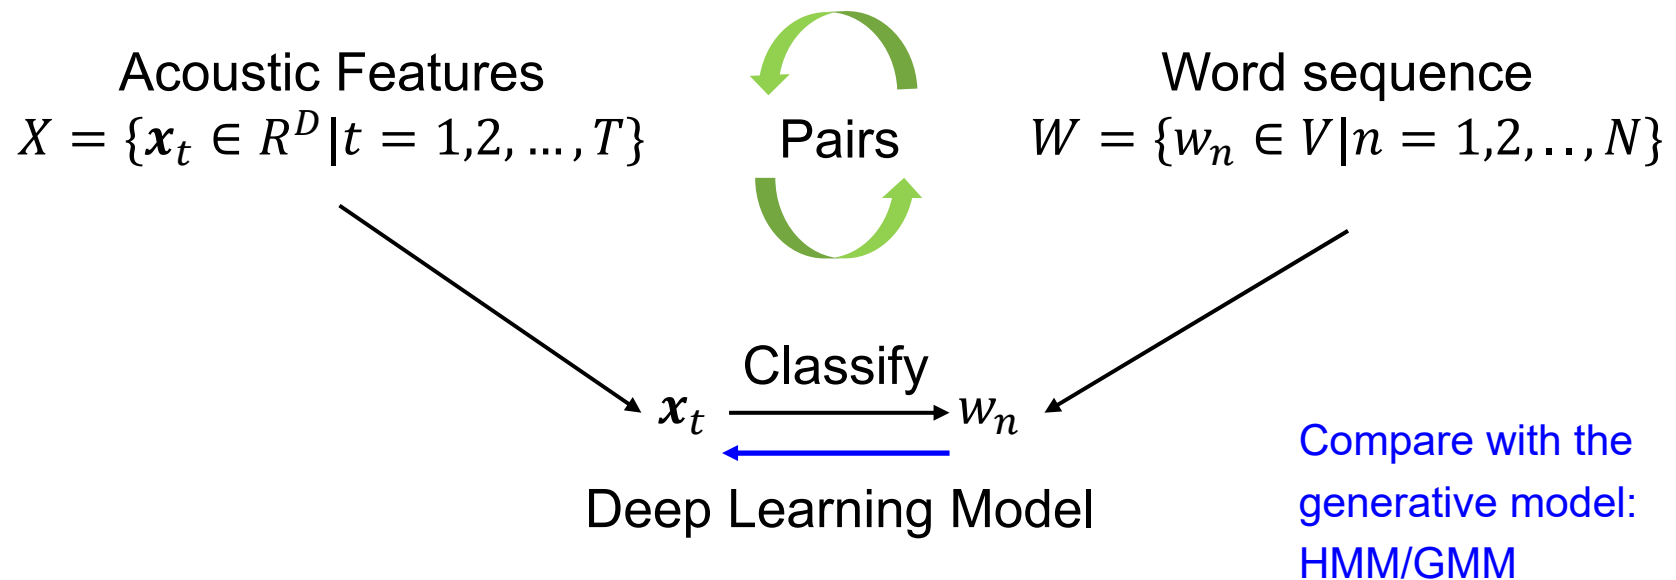
Part B: Development of ASR System

 Part C: End-to-End ASR System

Part D: Automatic Lyric Transcription (ALT)

# Overview of End-to-End ASR System

Modeling ASR task as a sequence-to-sequence classification problem



Problems:

- $X$  and  $W$  can vary in length, normally  $T \gg N$
- The ratio of the lengths of  $X$  and  $W$  can vary,  $\frac{T}{N}$  is different for each pair
- **Accurate alignment** of  $X$  and  $W$  is absent

# Overview of End-to-End ASR System

## Problems:


- $X$  and  $W$  can vary in length, normally  $T \gg N$
- The ratio of the lengths of  $X$  and  $W$  can vary,  $\frac{T}{N}$  is different for each pair
- Accurate alignment of  $X$  and  $W$  is absent

## Solutions:

- ***Connectionist Temporal Classification (CTC)***
- Encoder/Decoder with Attention
- Hybrid CTC-Attention
- Transducer



# Connectionist Temporal Classification (CTC)

- CTC is introduced to label unsegmented sequences directly
- CTC adopts an encoder to predict frame-level text transcript, or called alignment



# Motivation of CTC

- Given acoustic features  $X = [x_1, x_2, \dots, x_T]$  and output text transcription  $W = [w_1, w_2, \dots, w_N]$ , how can we align them?
- Naïve Method: remove the repeated tokens



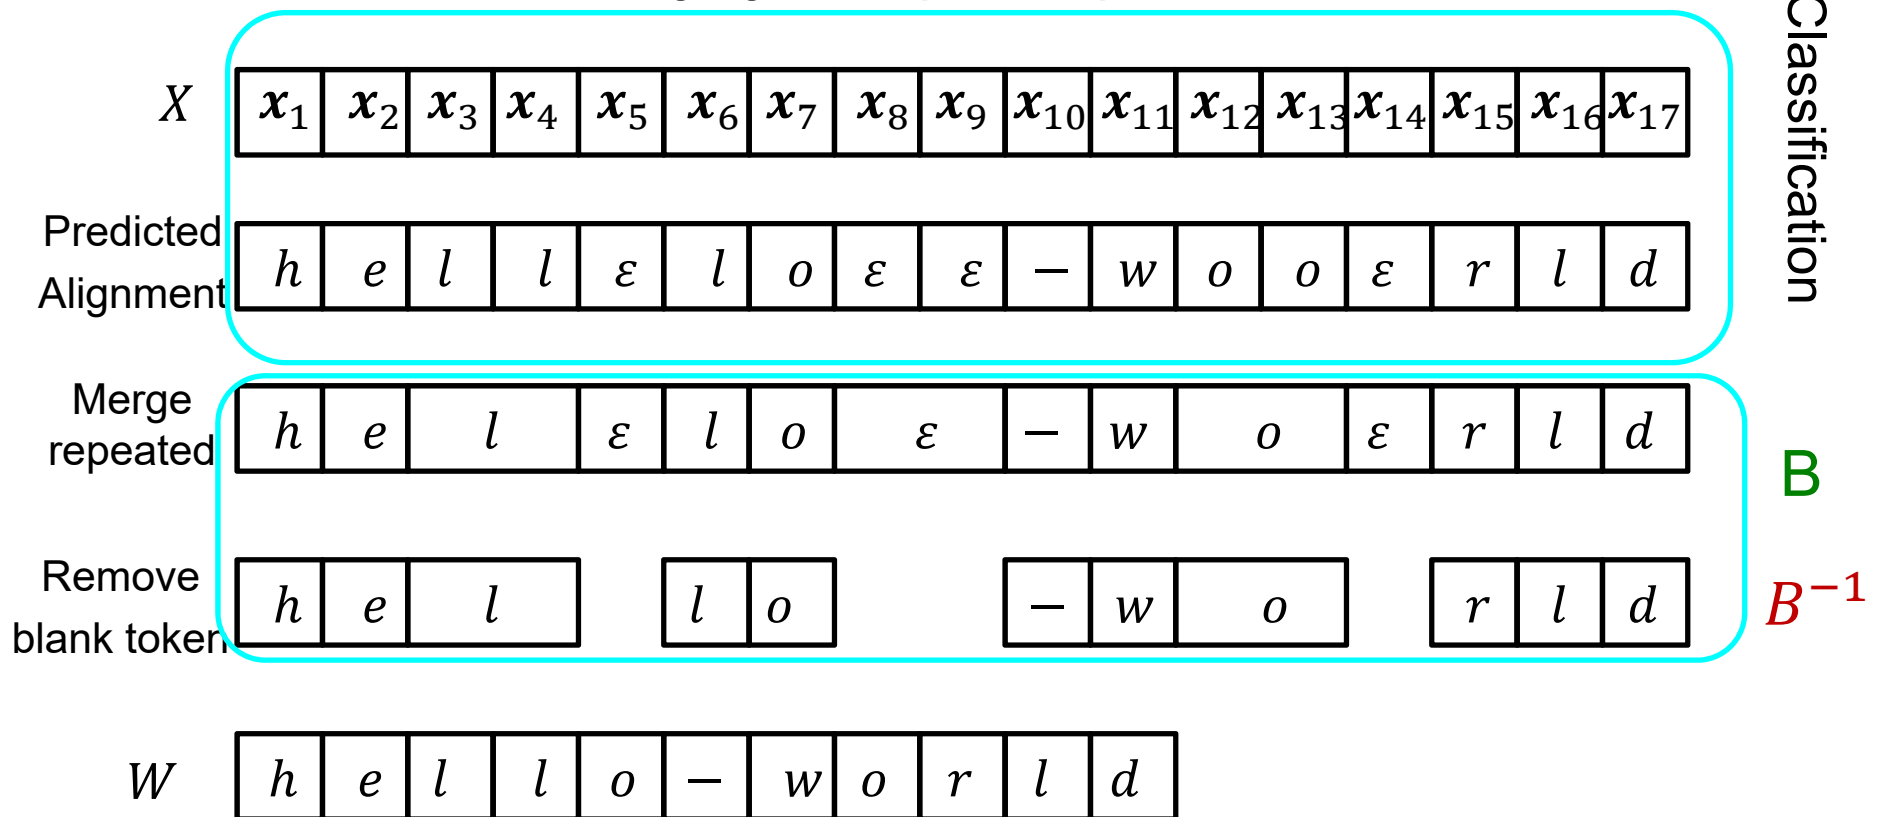
[1] Graves A, Fernández S, Gomez F, et al. Connectionist temporal classification: labelling unsegmented sequence data with recurrent neural networks[C]//Proceedings of the 23rd international conference on Machine learning. 2006: 369-376.

[2] Hannun A. Sequence modeling with ctc[J]. Distill, 2017, 2(11): e8.




# How does CTC work?

- CTC introduces a **blank token**  $\varepsilon$  (different from space token), which is removed after merging the repeated predicted tokens



# How to compute CTC loss?

- A correct text transcript corresponds to multiple alignments



- The probability  $P(W|X)$  can be represented as

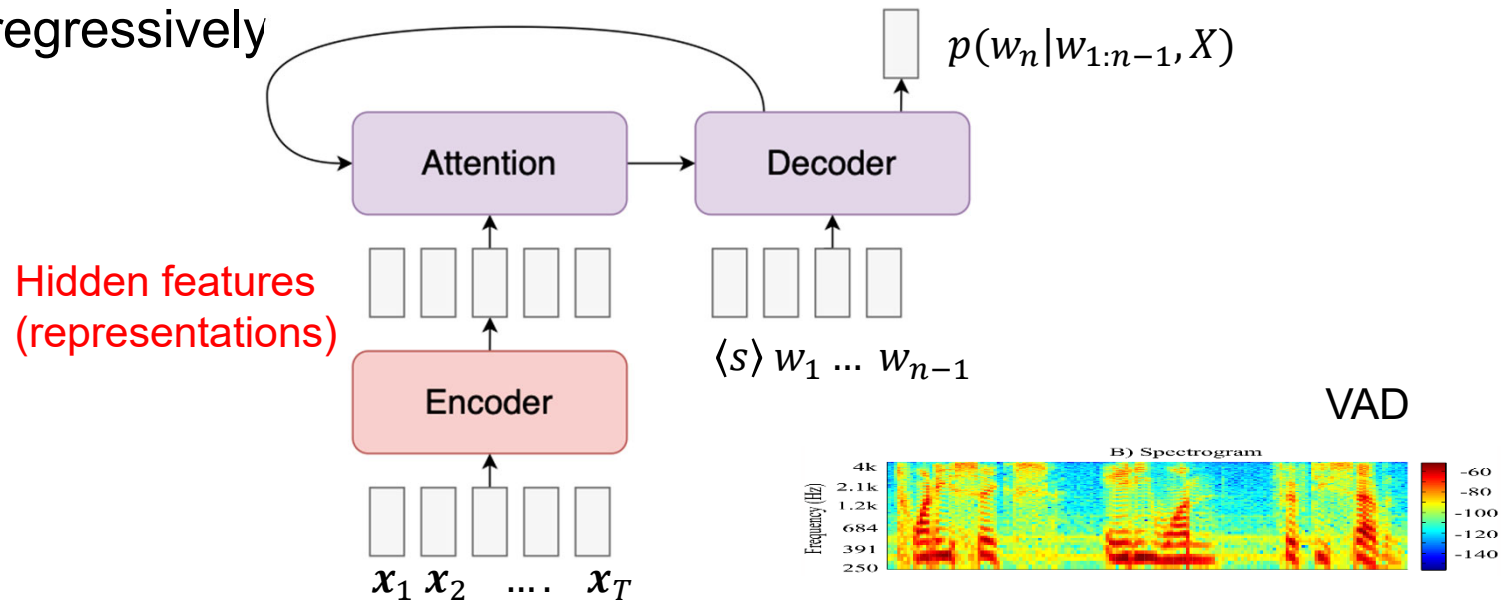
$$P(W|X) = \sum_{\pi \in B^{-1}(W)} \prod_{t=1}^T p(\pi_t | \mathbf{x}_t)$$

- Loss function can be written as:

$$L_{CTC} = -\log \sum_{\pi \in B^{-1}(W)} \prod_{t=1}^T p(\pi_t | \mathbf{x}_t)$$

# Encoder/Decoder with Attention

- Encoder: transform the acoustic features into a sequence of hidden features
- Attention: allows the decoder to pay attention to different parts of hidden features when predicting each token
- Decoder: accepts previous tokens and hidden features, decodes them autoregressively



[1] Bahdanau D, Chorowski J, Serdyuk D, et al. End-to-end attention-based large vocabulary speech recognition[C]//2016 IEEE international conference on acoustics, speech and signal processing (ICASSP). IEEE, 2016: 4945-4949.

# Encoder/Decoder with Attention

- Encoder: transform the acoustic features into a sequence of states
- Decoder: autoregressively predict the output tokens
- Attention: allows the decoder to pay attention to different parts of hidden states when predicting each token
- Attention type: content-aware attention / location-aware attention / multi-head attention

[1] Bahdanau D, Chorowski J, Serdyuk D, et al. End-to-end attention-based large vocabulary speech recognition[C]//2016 IEEE international conference on acoustics, speech and signal processing (ICASSP). IEEE, 2016: 4945-4949.

[2] Chorowski J K, Bahdanau D, Serdyuk D, et al. Attention-based models for speech recognition[J]. Advances in neural information processing systems, 2015, 28.

[3] Vaswani A, Shazeer N, Parmar N, et al. Attention is all you need[J]. Advances in neural information processing systems, 2017, 30.

# Encoder/Decoder with Attention

- The probability  $P(W|X)$  can be represented as

$$P(W|X) = \prod_{n=1}^N p(w_n|w_{1:n-1}, X)$$

- Loss function can be written as:

$$L_{S2S} = -\log \prod_{n=1}^N p(w_n|w_{1:n-1}, X)$$

# Hybrid CTC-Attention

- Multi-task learning: combining CTC and Encoder/Decoder Attention
- Training Loss function

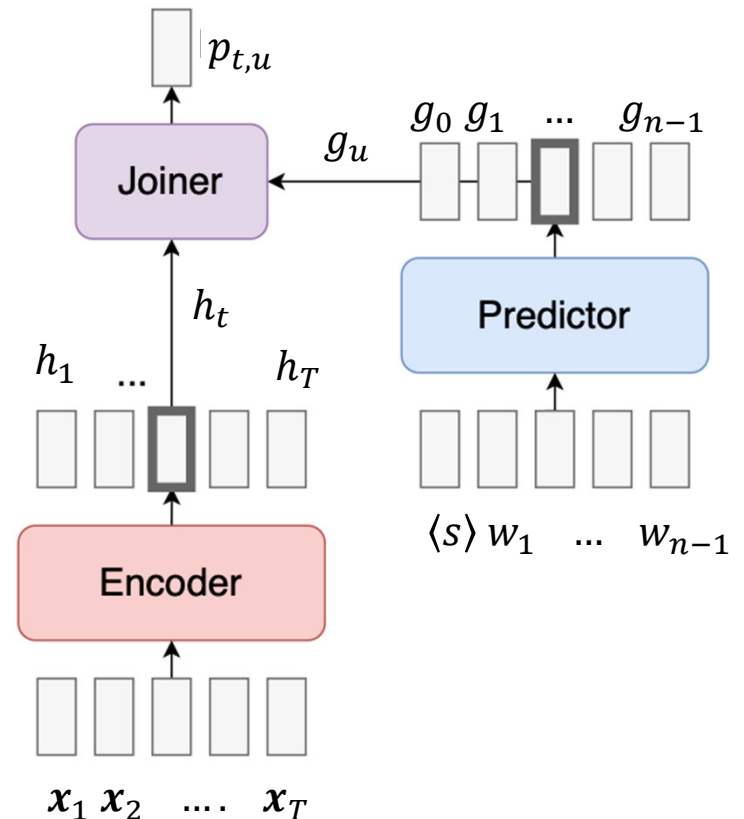
$$\begin{aligned} L_{Hybrid} &= \lambda L_{CTC} + (1 - \lambda) L_{S2S} \\ &= -\lambda \log \sum_{\pi \in B^{-1}(W)} \prod_{t=1}^T p(\pi_t | \mathbf{x}_t) \\ &\quad - (1 - \lambda) \log \prod_{n=1}^N p(w_n | w_{1:n-1}, X) \end{aligned}$$

[1] Kim S, Hori T, Watanabe S. Joint CTC-attention based end-to-end speech recognition using multi-task learning[C]//2017 IEEE international conference on acoustics, speech and signal processing (ICASSP). IEEE, 2017: 4835-4839.

# Transducer

- Encoder: transform the acoustic features into a sequence of hidden features
- Predictor: autoregressively predict the output tokens by only taking previous tokens
- Joiner: combine the hidden features and predictor outputs and outputs the distribution over possible tokens (including a  $\emptyset$  token)

$$p_{t,u} = P(w_{t,n} | w_{1:n-1}, X)$$



[1] Lugosch, Loren. "Sequence-to-sequence learning with Transducers", 2020.  
<https://lorenlugosch.github.io/posts/2020/11/transducer/>

# How to Implement an End-to-End ASR system?

Encoder, Decoder, Predictor, Joiner are implemented by deep learning models:

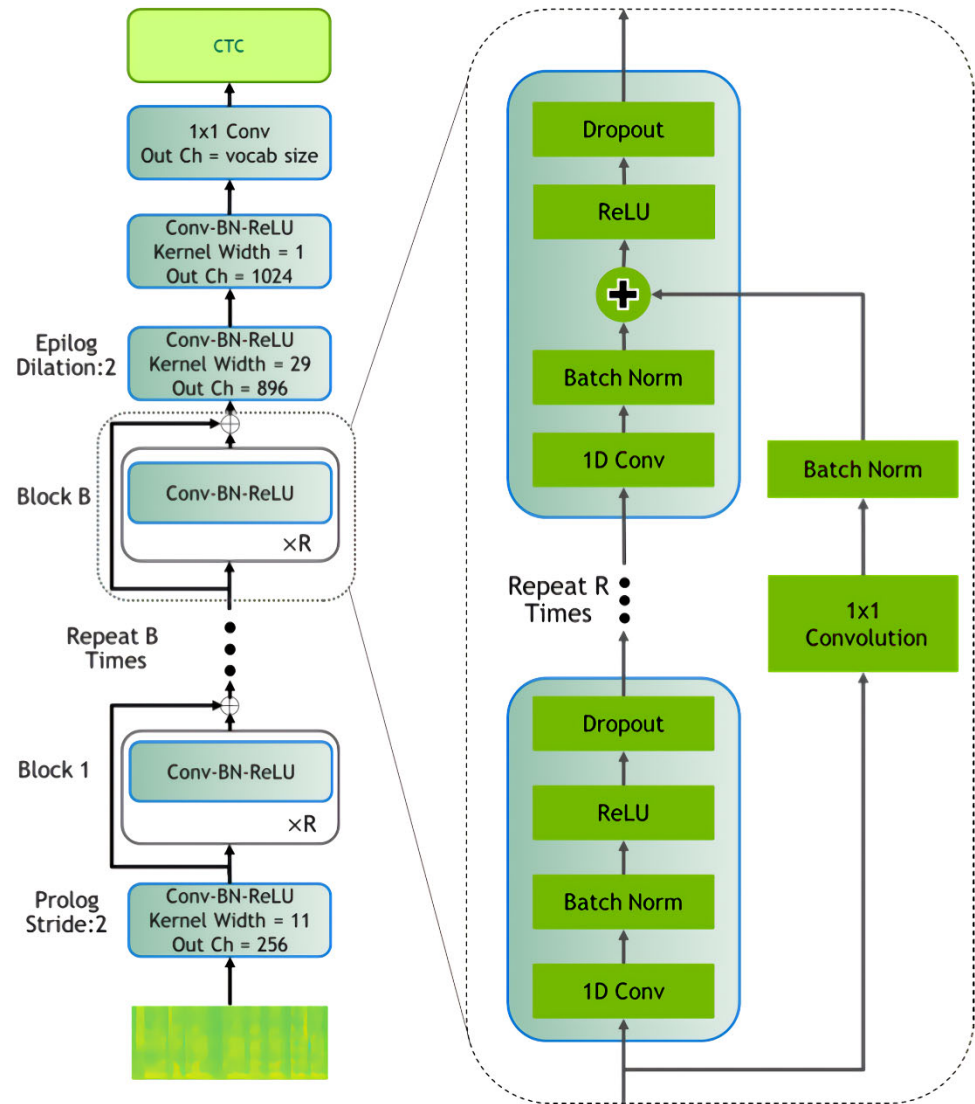
- Multi-layer Perceptron (MLP)
- Convolutional Neural Networks (CNN)
- Recurrent Neural Networks (RNN)
- Transformer
- Representative variants: Jasper, Conformer, wav2vec 2.0, etc.

Encoder performs representation learning, which makes it an important part of the whole model



# Example 1: Jasper

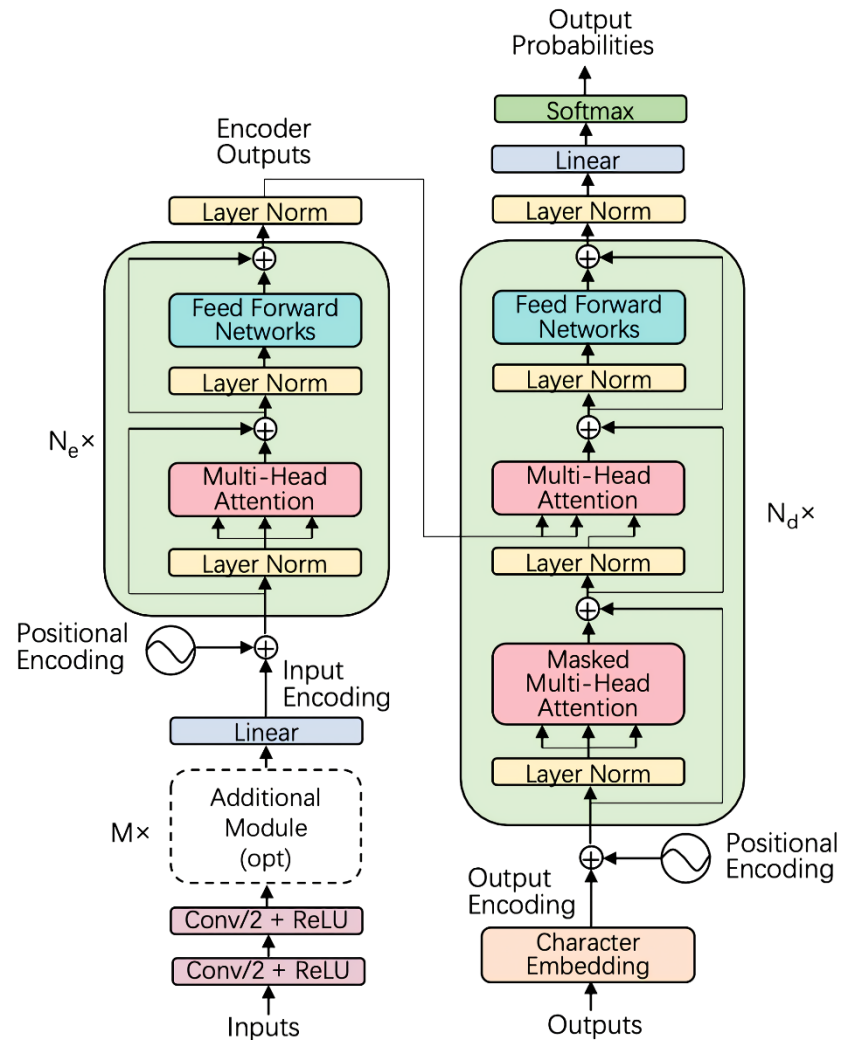
- Implemented the CTC encoder as a 1D-CNN
- Residual blocks and Dense blocks design



[1] Li J, Lavrukhin V, Ginsburg B, et al. Jasper: An end-to-end convolutional neural acoustic model[J]. arXiv preprint arXiv:1904.03288, 2019.

# Example 2: Speech-Transformer

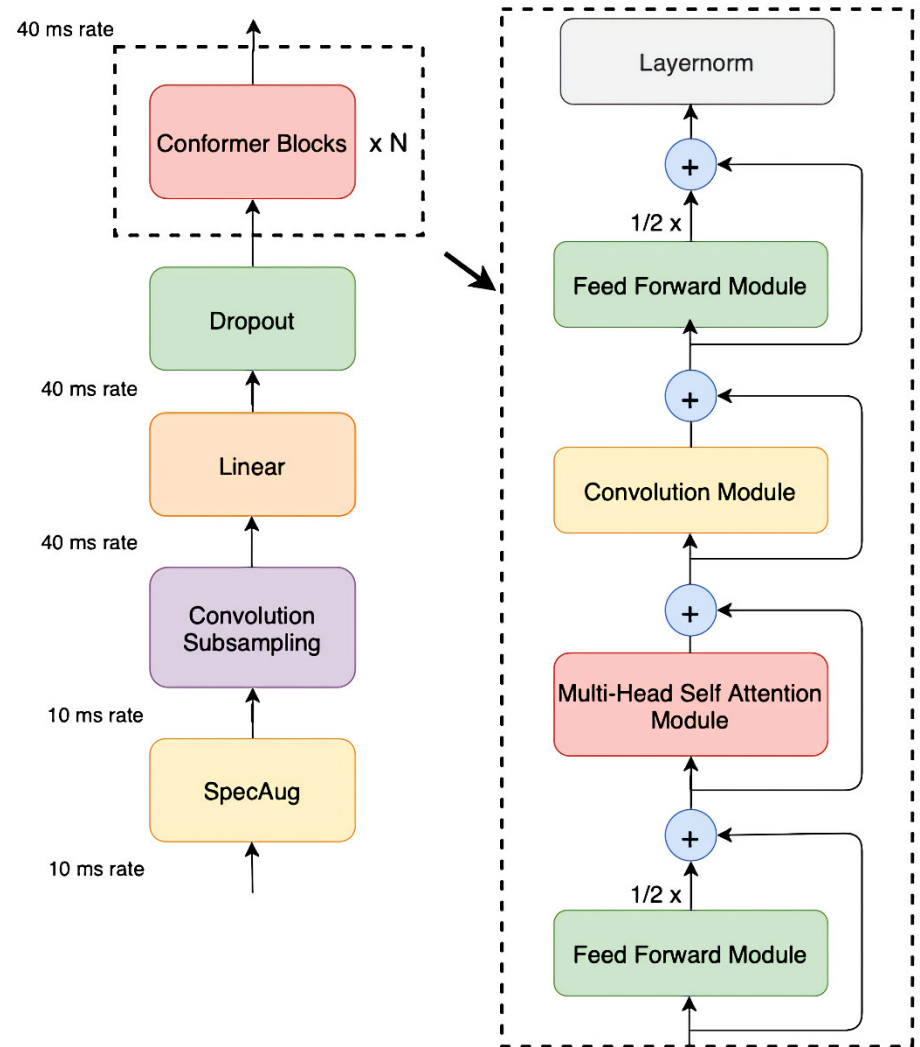
- Motivated by the success of Transformer in NLP field
- CNN are used to exploit the structure locality of inputs



[1] Dong L, Xu S, Xu B. Speech-transformer: a no-recurrence sequence-to-sequence model for speech recognition[C]//2018 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP). IEEE, 2018: 5884-5888.

# Example 3: Conformer


- Conformer combines both CNN and Transformer
- Fully exploit local/global dependencies among features of different frames



[1] Gulati A, Qin J, Chiu C C, et al. Conformer: Convolution-augmented transformer for speech recognition[J]. arXiv preprint arXiv:2005.08100, 2020.

# Example 4: wav2vec 2.0 architecture

- wav2vec 2.0 can be adopted as the encoder to extract deep representations from raw audio
- wav2vec 2.0 includes a CNN and a Transformer




[1] Baevski A, Zhou Y, Mohamed A, et al. wav2vec 2.0: A framework for self-supervised learning of speech representations[J]. Advances in Neural Information Processing Systems, 2020, 33: 12449-12460.

# Example 4: wav2vec 2.0 training

➤ Training wav2vec 2.0 has two stages:

Stage I: Self-supervised Contrastive Learning

Stage II: Supervised Fine-tuning




# Language Model (LM)

- HMM-based ASR: LM participates in the testing phase

$$\hat{W} = \operatorname{argmax}_{W \in V} P_{\Phi}(X|W) P_{\Theta}(W)$$

Language Model



- End-to-End ASR: LM can serve as **regularization** to original optimization objective, empirically improve ASR performance

$$\hat{W} = \operatorname{argmax}_{W \in V} P_{\Phi}(W|X)$$



$$\hat{W} = \operatorname{argmax}_{W \in V} P_{\Phi}(W|X) P_{\Theta}(W)^{\beta}$$

Language Model



# What is LM?

- LM is a system which can predict the next token  $w_{t+1}$  given a sequence of previous tokens  $w_1, \dots, w_t$

$$P(w_{t+1} | w_1, w_2, \dots, w_t)$$

- LM can also model the probability of a sequence of tokens  $w_1, \dots, w_T$

$$P(w_1, w_2, \dots, w_T) = \prod_{t=1}^T \underbrace{P(w_t | w_1, \dots, w_{t-1})}_{\text{LM's output}}$$

- LM families: ***n*-gram LM**, RNNLM, Transformer LM

# n-gram LM definition

- An *n-gram* is a chunk of n consecutive tokens, take words as examples

**Unigrams:** “he”, “is”, “a”, “hero”

**Bigrams:** “he is”, “is a”, “a hero”

**Trigrams:** “he is a”, “is a hero”

**4-grams:** “he is a hero”

- *(n-1)*-order **Markov assumption**: *n-gram* LM assumes  $w_{t+1}$  only depends on the previous *n-1* tokens  $w_{t-n+2}, \dots, w_t$

$$P(w_{t+1}|w_1, w_2, \dots, w_t) = P(w_{t+1}|w_{t-n+2}, \dots, w_t)$$

$$= \frac{P(w_{t-n+2}, \dots, w_t, w_{t+1})}{P(w_{t-n+2}, \dots, w_t)}$$

n-gram prob

(n-1)-gram prob



# How to build an n-gram LM?

- How to compute n-gram probability and (n-1)-gram probability?

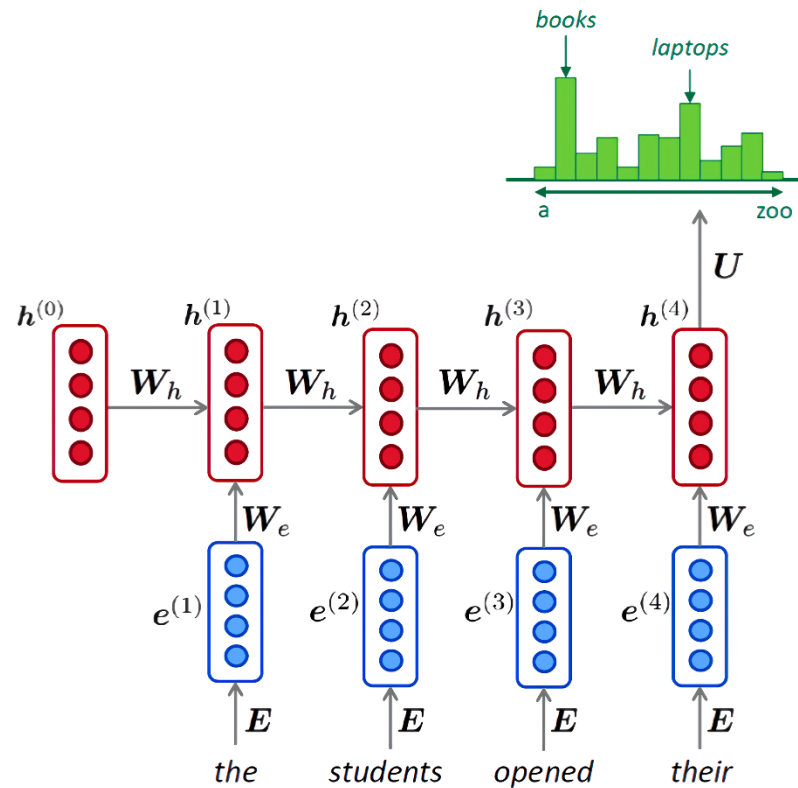
$$P(w_{t+1}|w_1, w_2, \dots, w_t) = \frac{P(w_{t-n+2}, \dots, w_t, w_{t+1})}{P(w_{t-n+2}, \dots, w_t)}$$

n-gram prob  
(n-1)-gram prob

- Solution: Count the frequencies of n-grams and (n-1)-grams in ***large corpus of text***
- n-gram LM is a probabilistic model without deep learning

# RNN LM definition

- Directly model the  $P(w_{t+1}|w_1, w_2, \dots, w_t)$  by deep neural networks
- To enable the inputs with arbitrary length, RNN is adopted



E.g., Given four tokens, predict the next token

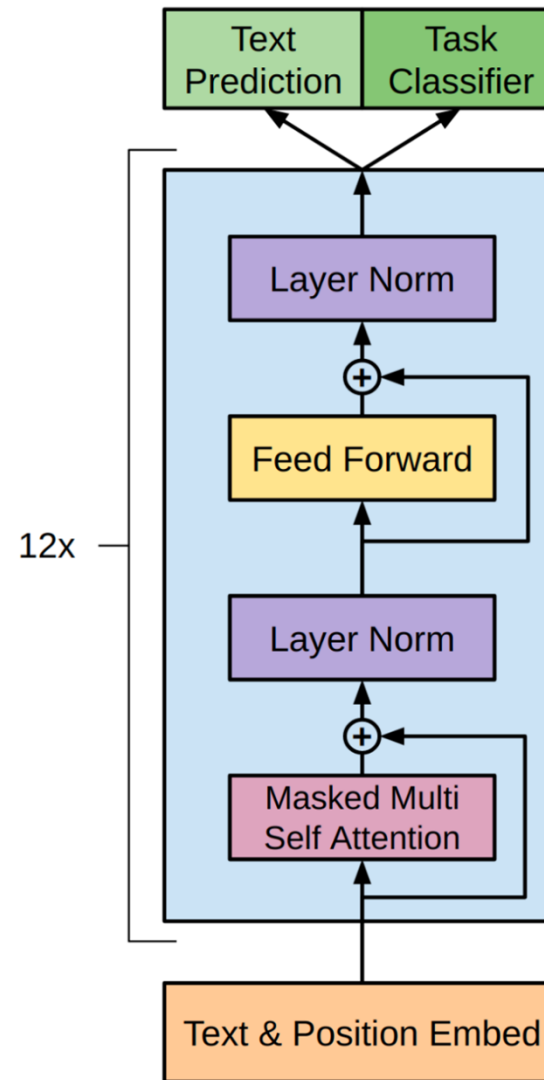
# RNNLM implementation

- Directly model the  $P(w_{t+1}|w_1, w_2, \dots, w_t)$  by deep neural networks
  - To enable the inputs with arbitrary length, RNN is adopted
  - LSTM and GRU can alleviate vanishing gradient problem comparing to vanilla RNN
- 
- Train an RNNLM: loss function is defined as cross-entropy between predicted probability distribution of next token and true next token
  - Evaluate an RNNLM: standard evaluation metric is perplexity

$$\text{perplexity} = \prod_{t=1}^T \left( \frac{1}{P_{LM}(w_{t+1}|w_1, w_2, \dots, w_t)} \right)^{1/T}$$

# Transformer LM

- Model  $P(w_{t+1}|w_1, w_2, \dots, w_t)$  by Transformer
- Transformer LM is capable of modeling long dependencies among tokens
- Training and evaluation of Transformer LM are like RNNLM



[1] Vaswani A, Shazeer N, Parmar N, et al. Attention is all you need. Advances in neural information processing systems, 2017, 30.

# Decoding for ASR

- HMM-based ASR:

$$\hat{W} = \operatorname{argmax}_{W \in V} P_{\Phi}(X|W)P_{\Theta}(W) = \operatorname{argmax}_{W \in V} \log P_{\Phi}(X|W) + \log P_{\Theta}(W)$$


- End-to-End ASR:

$$\hat{W} = \operatorname{argmax}_{W \in V} P_{\Phi}(W|X)P_{\Theta}(W)^{\beta} = \operatorname{argmax}_{W \in V} \log P_{\Phi}(W|X) + \beta \log P_{\Theta}(W)$$

- But how to optimize the above objectives, thus searching for the optimal  $W$ ?


# CTC Decoding

Take the CTC decoding for example. Suppose we use character as the token,  $X$  has  $T$  frames and the vocabulary has  $C$  tokens




# Brute Force

- Check all the possible path, combine all identical tokens (merge repeated tokens and remove blank token) and compare the probability of the token sequences
- Computational complexity is  $O(C^T)$ , best performance



# Best path Decoding


- Choose the most likely character per frame
- May not be the most probable token sequence (text transcript) because a text transcript corresponds to multiple paths
- Computational complexity is  $O(TC)$ , suboptimal performance






# Beam Search

- At each frame, choose the ***B text candidates (beams)*** with highest probability and store them for the next frame
- NOTE: beams are the **text candidates not the alignment candidates**




# Beam Search

- At each frame, choose the ***B text candidates (beams)*** with highest probability and store them for the next frame
- NOTE: beams are the ***text candidates not the alignment candidates***




# Beam Search

- At each frame, choose the ***B text candidates (beams)*** with highest probability and store them for the next frame
- NOTE: beams are the ***text candidates not the alignment candidates***




# Beam Search

- At each frame, choose the ***B text candidates (beams)*** with highest probability and store them for the next frame
- NOTE: beams are the ***text candidates not the alignment candidates***




# Beam Search

- At each frame, choose the ***B text candidates (beams)*** with highest probability and store them for the next frame
- NOTE: beams are the ***text candidates not the alignment candidates***




# Beam Search

- At each frame, choose the ***B text candidates (beams)*** with highest probability and store them for the next frame
- NOTE: beams are the ***text candidates not the alignment candidates***




# Beam Search

- At each frame, choose the ***B text candidates (beams)*** with highest probability and store them for the next frame
- NOTE: beams are the ***text candidates not the alignment candidates***



# Beam Search


- At each frame, choose the ***B text candidates (beams)*** with highest probability and store them for the next frame
- NOTE: beams are the ***text candidates not the alignment candidates***






# Beam Search

- Finally, choose the best beam as the prediction
- Computational Complexity:  $O(TCB)$



# Beam Search

- Please refer to the following material for detailed explanation  
<https://towardsdatascience.com/beam-search-decoding-in-ctc-trained-neural-networks-5a889a3d85a7>



# Topics Today

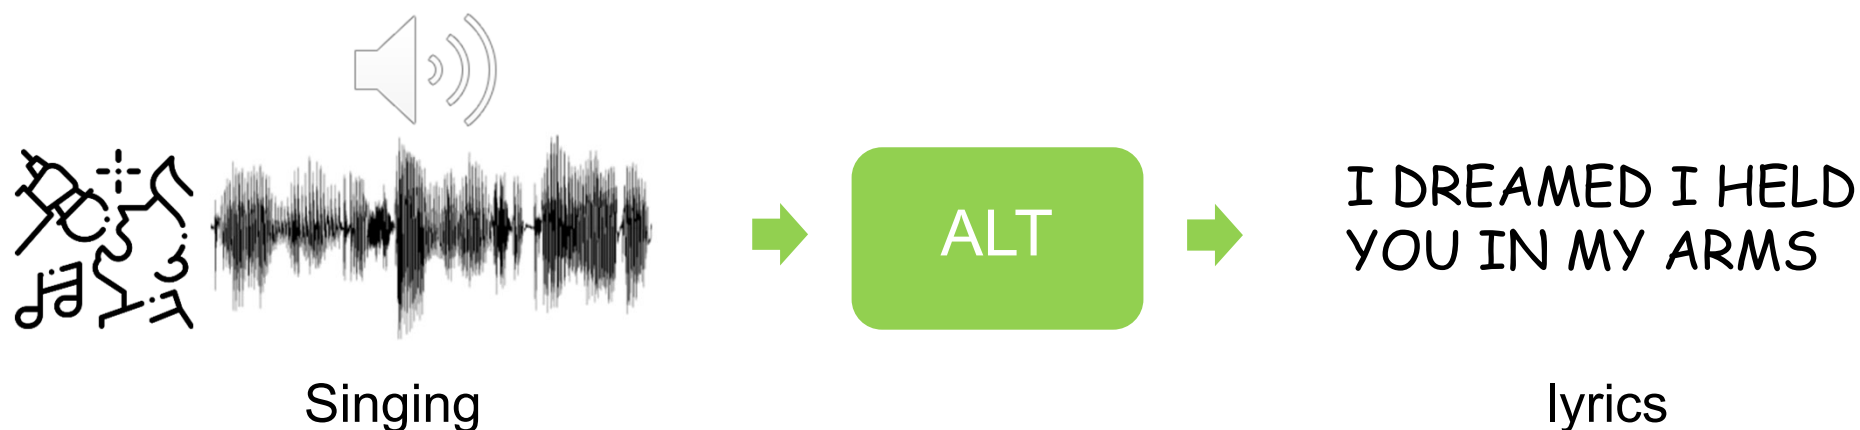
Part A: Overview of Automatic Speech Recognition (ASR)

Part B: Development of ASR System

Part C: End-to-End ASR System

 Part D: Automatic Lyric Transcription (ALT)

# What is Automatic Lyric Transcription (ALT)?



- ALT task aims to transcribe a singing waveform into lyrics
- A counterpart task of ASR

# Relationship between ALT and ASR

## Similarity:

- Both ASR and ALT aim to transcript the text from audio
- Singing voice and speech are produced by the same organ
- The target transcription have the same vocabularies

## Discrepancies:

- Singing voice is less intelligible and harder to be recognized
- Pitch and duration are different

[1] Zhang C, Yu J, Chang L C, et al. PDAugment: Data Augmentation by Pitch and Duration Adjustments for Automatic Lyrics Transcription[J]. arXiv preprint arXiv:2109.07940, 2021.

[2] Gu X, Ou L, Ong D, Wang Y. MM-ALT: A Multimodal Automatic Lyric Transcription System[C]//Proceedings of the 30th ACM International Conference on Multimedia. 2022.

# How to build an ALT system?

Take advantage of similarities between ASR and ALT


- Adopt the same pipeline to train an ALT system

How to deal with the discrepancies

- Adjust speech data to generate “song-like” data
- Adapt the knowledge from ASR model to ALT model

# Example 1: PDAugment


- Generate “song-like” training data by adjusting ***duration*** of speech



[1] Zhang C, Yu J, Chang L C, et al. PDAugment: Data Augmentation by Pitch and Duration Adjustments for Automatic Lyrics Transcription[J]. arXiv preprint arXiv:2109.07940, 2021.

# Example 1: PDAugment

- Generate “song-like” training data by adjusting ***pitch*** of speech

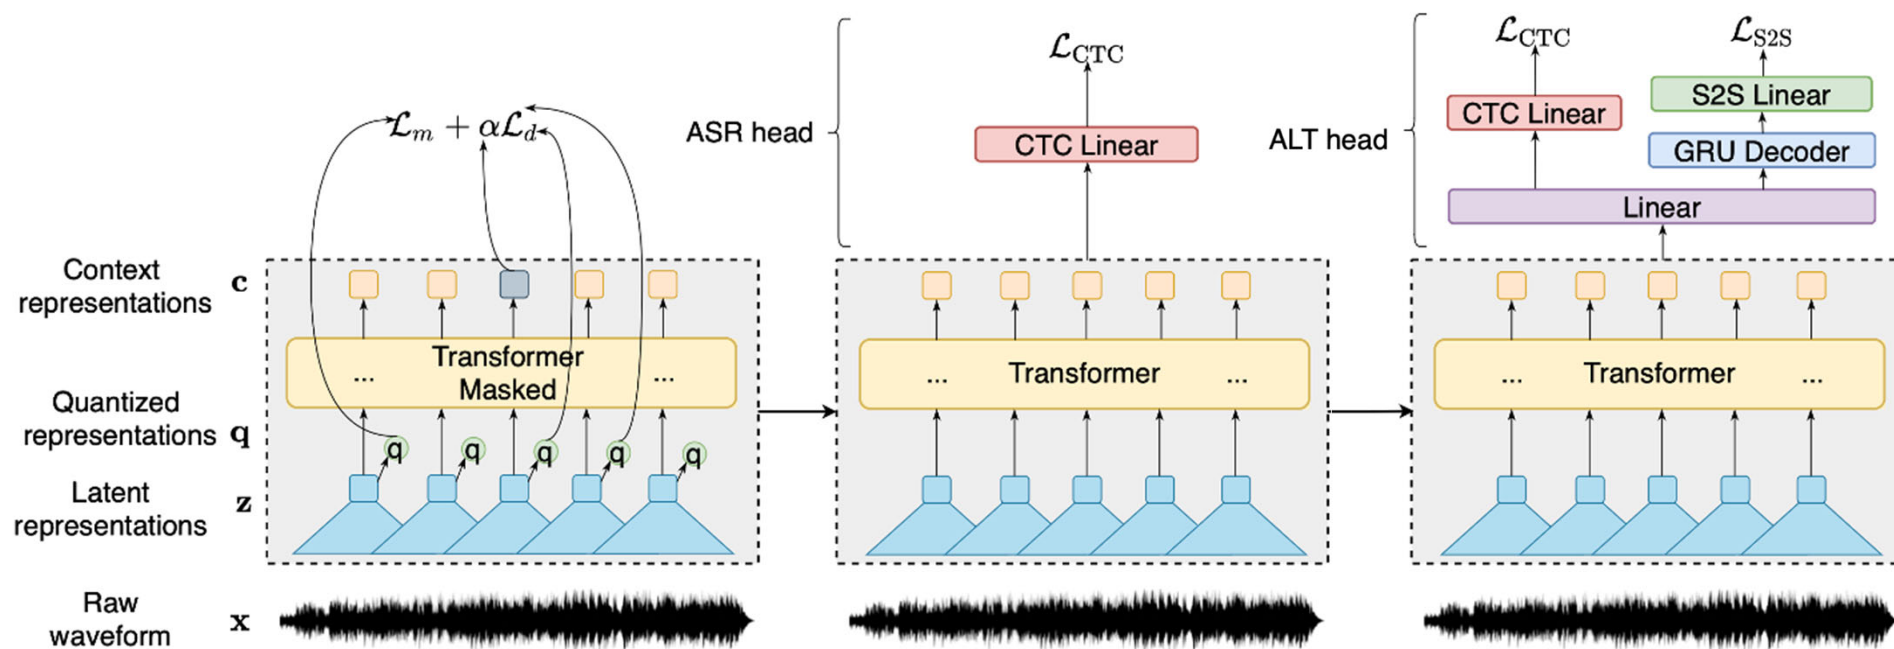


[1] Zhang C, Yu J, Chang L C, et al. PDAugment: Data Augmentation by Pitch and Duration Adjustments for Automatic Lyrics Transcription[J]. arXiv preprint arXiv:2109.07940, 2021.



## Example 2: Transfer Learning

- Fetch a wav2vec 2.0 after pre-training and fine-tuning on speech data
- Modify the head network and finetune the model on singing data
- Inherit the low-resource learning property



[1] Gu X, Ou L, Ong D, Wang Y. MM-ALT: A Multimodal Automatic Lyric Transcription System[C]//Proceedings of the 30th ACM International Conference on Multimedia. 2022.

# Take home message

To implement ASR and ALT systems, please consider the following platforms:

- SpeechBrain: <https://github.com/speechbrain/speechbrain> (recommended)
- Fairseq: <https://github.com/facebookresearch/fairseq> (more powerful but more challenging)

If interested, please also consider the following material:

- Tutorial on SpeechBrain: [https://colab.research.google.com/drive/1aFgZrUv3udM\\_gNJNUoLaHIm78QHtxdlz?usp=sharing#scrollTo=gzNb1cZzvxUh](https://colab.research.google.com/drive/1aFgZrUv3udM_gNJNUoLaHIm78QHtxdlz?usp=sharing#scrollTo=gzNb1cZzvxUh)
- Stanford lecture: <https://www.youtube.com/watch?v=3MjlkWxXigM>

Appendix 1: A nice introduction to Bayes  
Theorem, HMM and Viterbi Decoding.  
You are encouraged to watch it!

A friendly introduction to Bayes Theorem and Hidden Markov Models

By Luis Serrano


<https://www.youtube.com/watch?v=kqSzLo9fenk>

## Appendix 2: End-to-End Models for Speech Processing

<https://www.youtube.com/watch?v=3MjlkWxXigM>

### Speech Recognition -- the classical way

- Inference: Given audio features  $\mathbf{X} = x_1 x_2 \dots x_T$  infer most likely text sequence  $\mathbf{Y}^* = y_1 y_2 \dots y_L$  that caused the audio features



**Audio**

Features

HMM-GMM

Dictionary

**Text**

Here is another nice tutorial on ASR:

<https://www.youtube.com/watch?v=q67z7PTGRi8>

## Appendix 3: An intuitive introduction to language model!

The A.I. Hacker - Michael Phi

I Built a Personal Speech Recognition System for my AI Assistant

<https://www.youtube.com/watch?v=Yerel6Gn3bM>

## Appendix 4: What is a language model then?

Language modelling is to assign probability values to sequences of words.  
A subject studied by *computational linguists*.

The probability values help us to answer the following questions, for example:

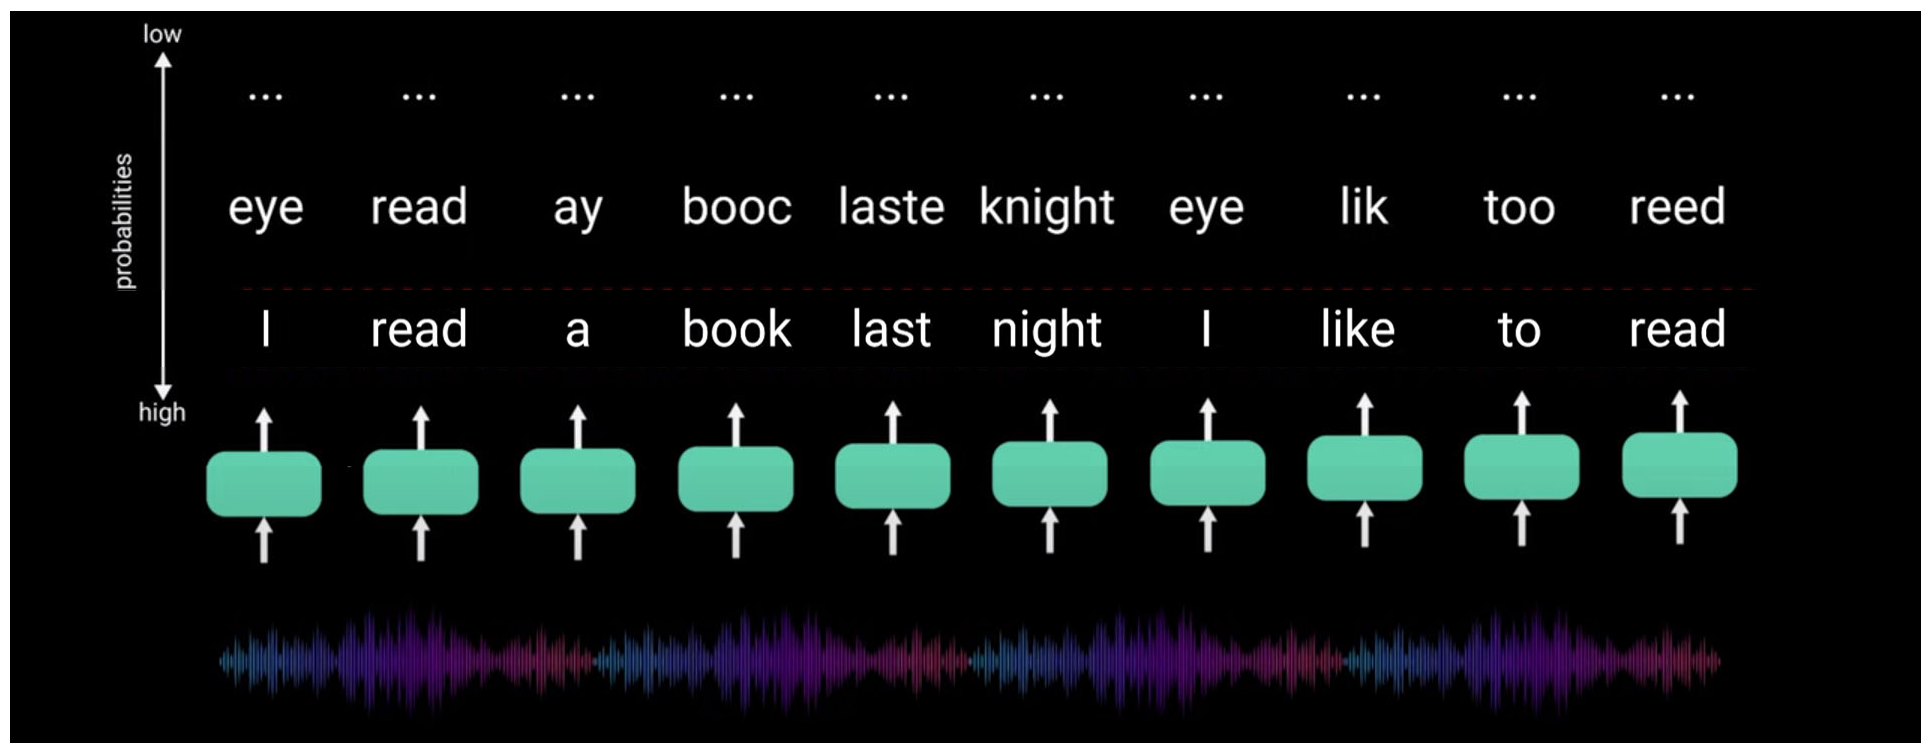
Which sentence is **grammatically** correct?

$$P(\text{"he eat pizza"}) < P(\text{"he eats pizza"})$$

Which **word order** is correct?

$$P(\text{"love I cats"}) < P(\text{"I love cats"})$$

## Appendix 5: If we only have an acoustic model



A pronunciation token can be different words!

## Appendix 6: Language model explained

Whats a more likely sentence?

Probability(I read a book) = 0.95

Probability(I red a book) = 0.25



## Appendix 7: Language model explained

### hypothesis (beams)

Probability( i read a book last night i like to read ) = 0.95

Probability( i red a book last night i like to read ) = 0.35

Probability( i red a book last night i like to reed ) = 0.15